

# Hardware Verification

## 2IMF20

Julien Schmaltz

Lecture 06:  
Verification of on-chip communication networks



Technische Universiteit  
**Eindhoven**  
University of Technology

Where innovation starts

# Living a revolution


---

“It is a really exiting time to be in computing right now, because we are [experiencing] a major change. A change that happens maybe once every 20 or 30 years ... a fundamental re-thinking of ... what computation is”

*William Dally, DAC keynote 2009*

# Multi-core shift reality


---






# Growing number of cores


Intel **8** cores  
~2.3 Bill. T. on 6.8cm<sup>2</sup>



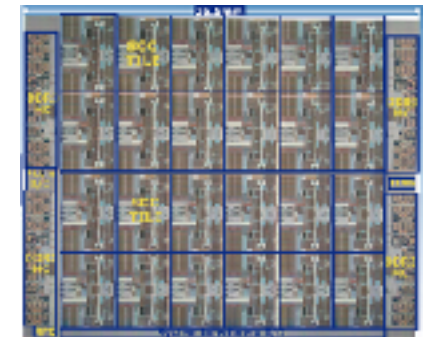
AMD Opteron **12** cores  
~1.8 Bill. T. on 2x3.46cm<sup>2</sup>




Sun Niagara3 **16** cores  
~1 Bill. T. on 3.7cm<sup>2</sup>



Intel SCC **48** cores  
~1.3 Bill. T. on 5.6cm<sup>2</sup>




Intel **4** cores  
~582 Mio. T. on 2.86cm<sup>2</sup>




Intel Research **80** cores  
~100 Mio. T. on 2.75cm<sup>2</sup>



Tilera TILEPro64 **64** cores



Intel **2** cores  
~167 Mio. T. on 1.1cm<sup>2</sup>



# 80 Cores Research Chip (Intel)

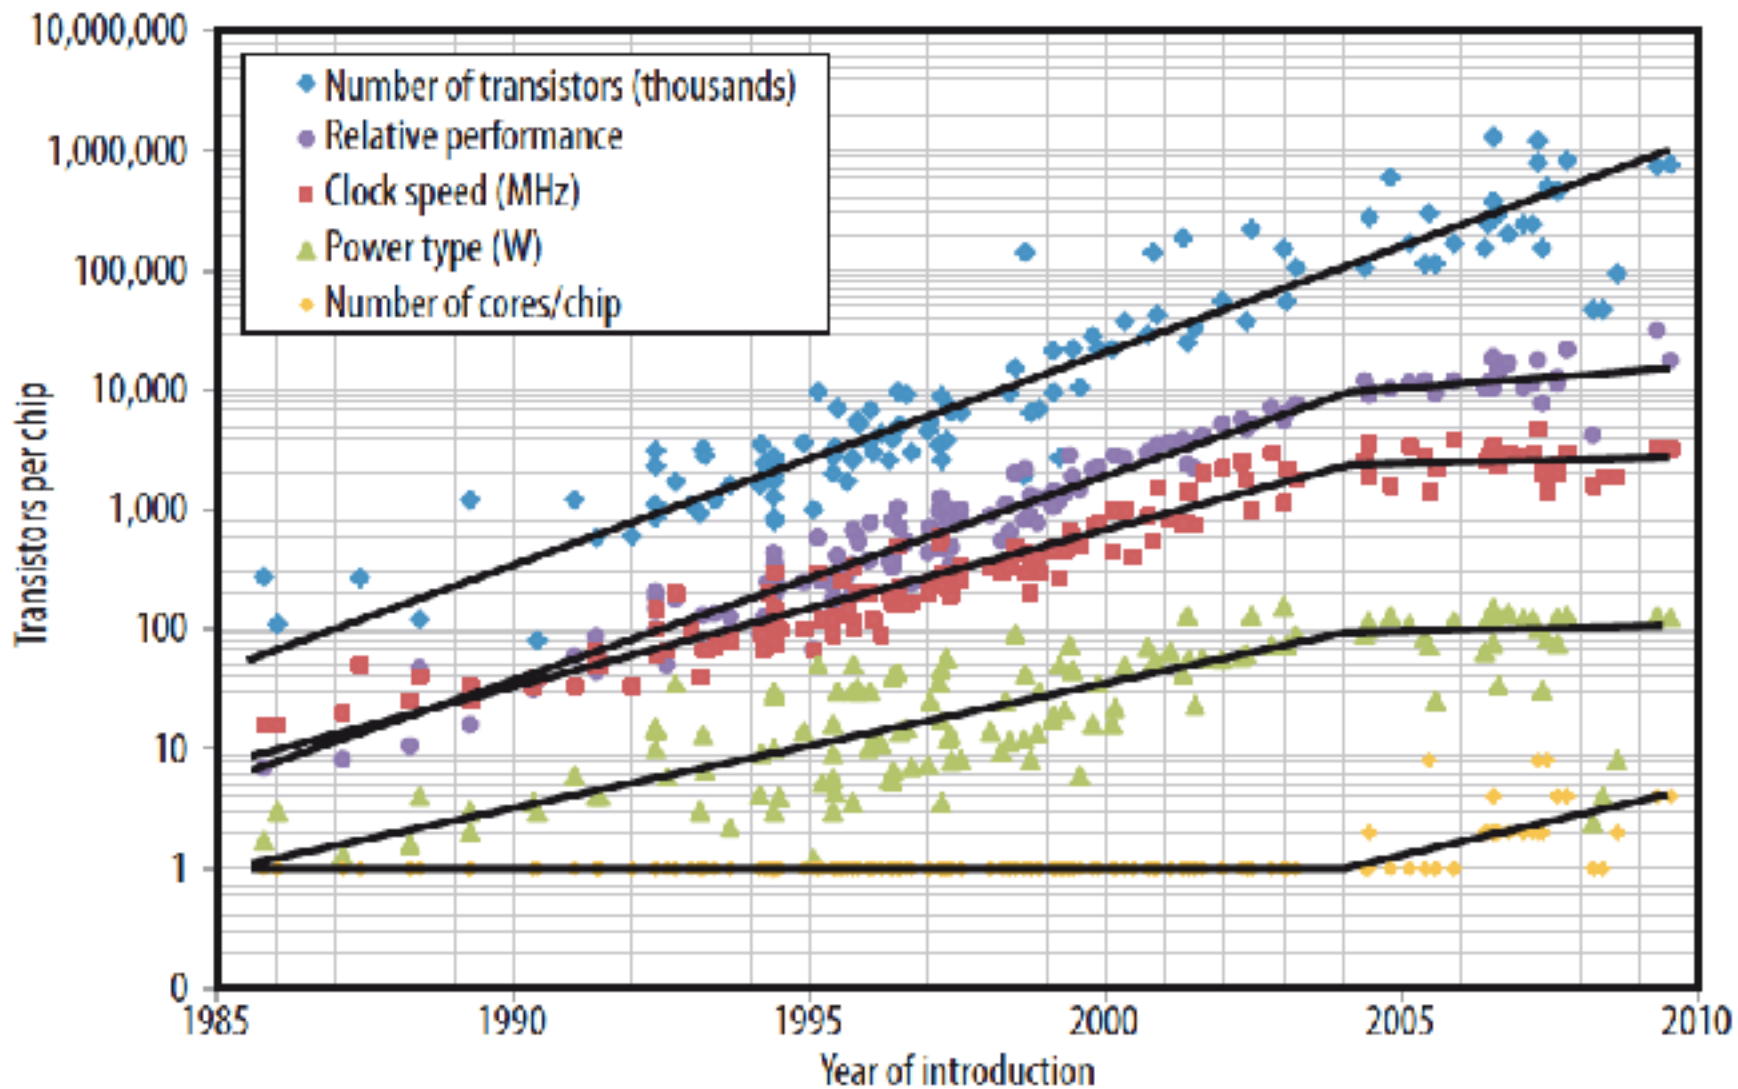
---

- Teraflops, 62 Watts
- 100 millions transistors, 275 mm<sup>2</sup>
- 25% node area for router



- ASCI Red Supercomputer
- Teraflops (Dec. 1996)
- 10, 000 Pentium Pro
- 104 cabinets, 230 m<sup>2</sup>

# It is only the beginning ...



**Figure 1.** Transistors, frequency, power, performance, and processor cores over time. The original Moore's law projection of increasing transistors per chip remains unabated even as performance has stalled.

Source. *IEEE Computers* 2011



# Communications are key

---


- 80 core research chip
  - about 25% area for the NoC
  - about 30% power consumption for the NoC
- Communication fabrics key
  - to performance and efficiency
  - to **functional correctness**

# Verification Challenges

---


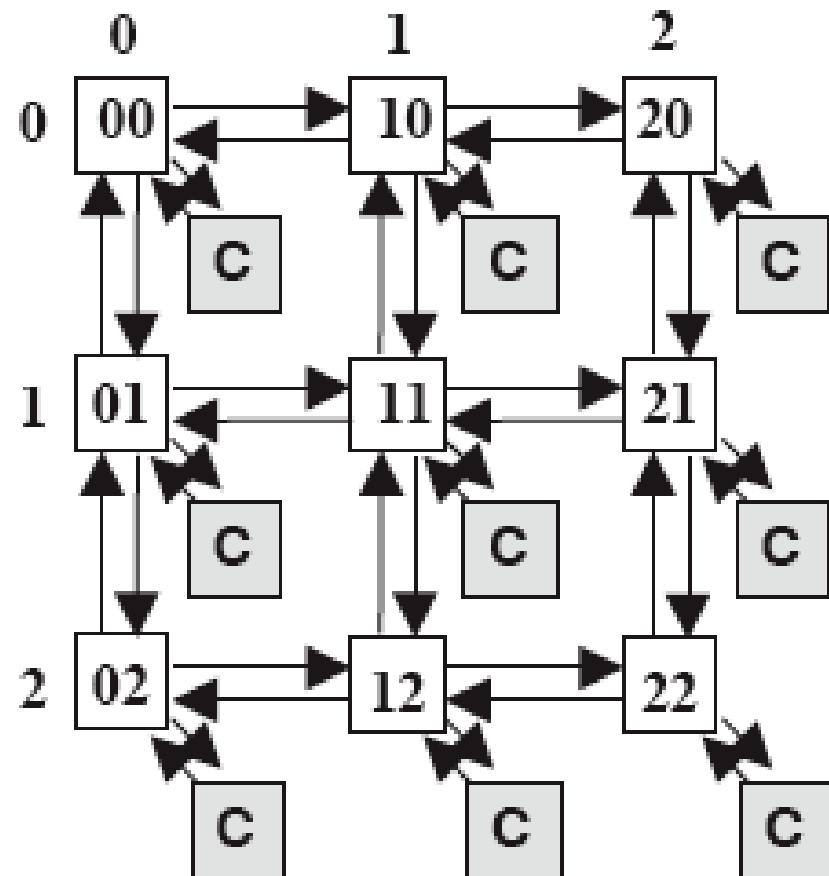
- NoCs are very large systems
  - Methods must scale up to **100s of agents**
  - Large number of parameters (routing, switching, buffers, etc.)
  - Regular and irregular structures
- NoCs must be fault-tolerant
  - Deep sub-micron effect
  - Not all routers/processors are working
  - Static and dynamic fault-models**
- NoCs have intricate message dependencies
  - Mix **between interconnect and protocols**
  - e.g. cache coherency or master/slave
  - Deadlocks can emerge from deadlock-free routing and protocols





# NoC Example: Hermes


---



# Application layer

---

- Masters send requests and wait for responses
- Slaves produce responses when receiving requests
- Deadlock-free protocol




No cyclic message dependencies:

# Network layer

---

- Deterministic simple routing algorithm
- First route to the destination column and then to the correct row
- No cyclic dependencies and thus deadlock-free





# Link layer

---


- A packet moves if the next channel has free space



- A packet moves if it has sufficiently credits



- A packet is joined with another packet



# Deadlock-free?

---

Deadlock-free  
application layer

+ Deadlock-free  
network layer

+ Deadlock-free  
link layer


?  
=

Deadlock-free system

# All masters right, slaves left

---


» Is the system deadlock-free ?



# All masters right, slaves left

---

- » Is the system deadlock-free ?
- » Yes ! A deadlock would require a response to wait for a request

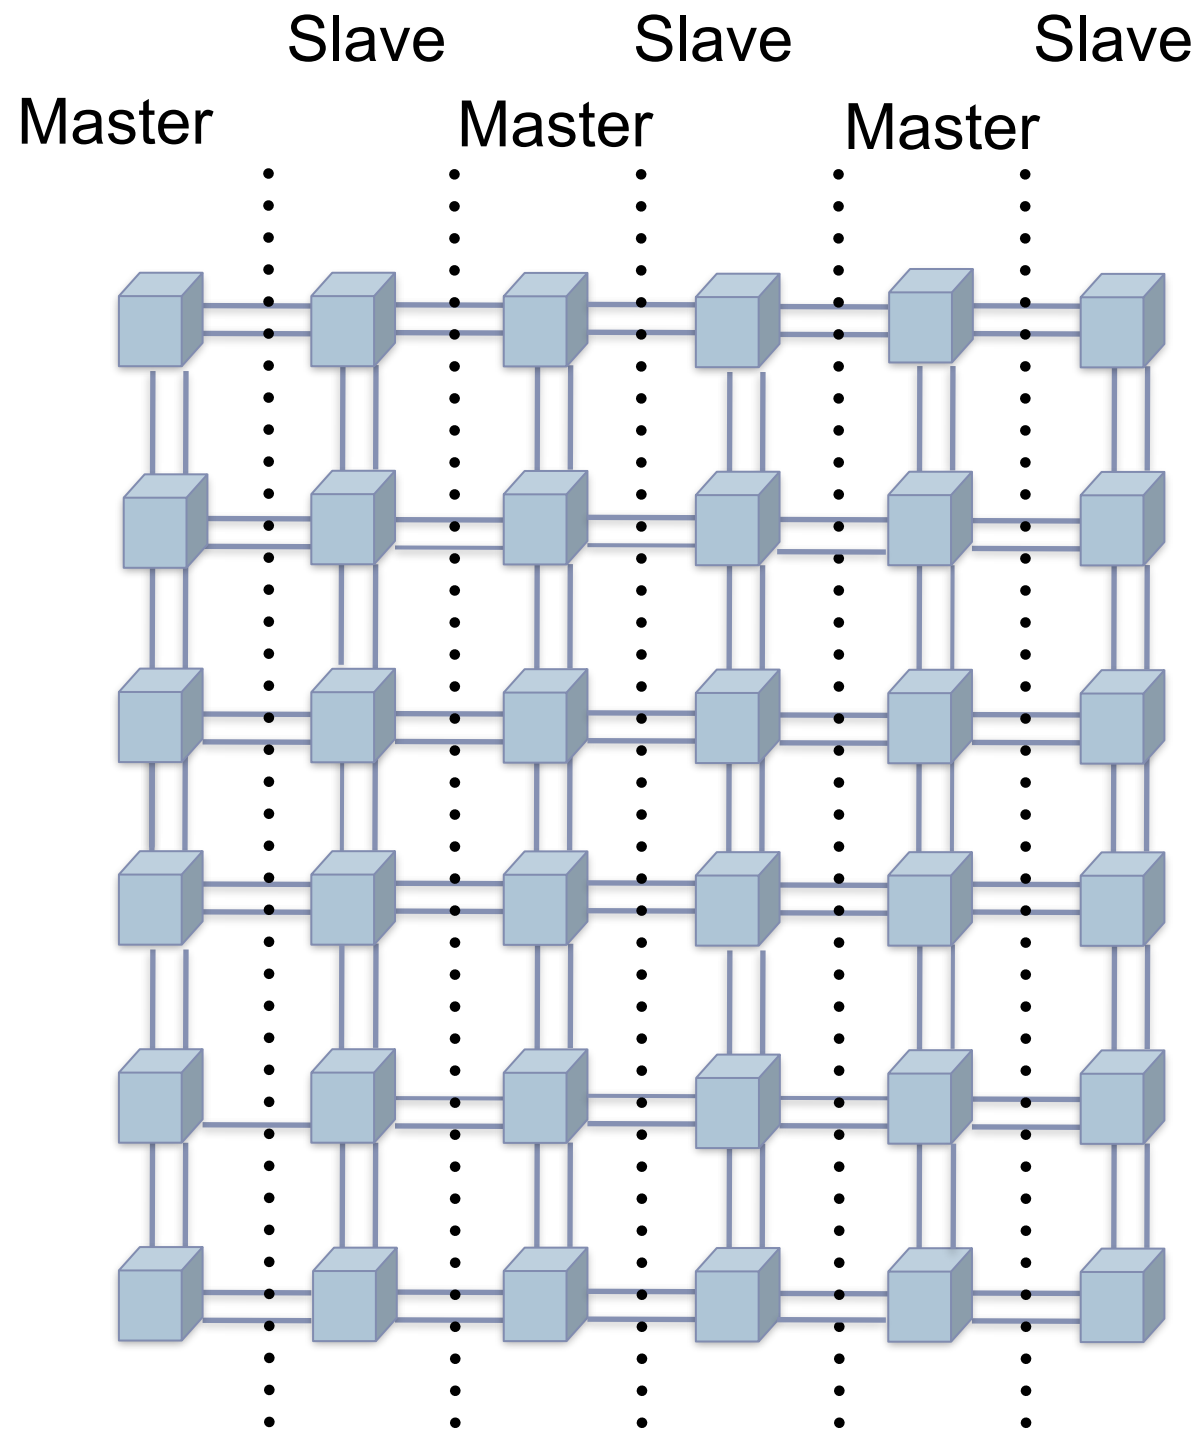




# Masters even columns, slaves odd


---

» Is the system deadlock-free ?




# Masters even columns, slaves odd

- » Is the system deadlock-free ?
- » No if at least four columns, yes otherwise.



# Masters even columns, slaves odd

- » Is the system deadlock-free ?
- » No if at least four columns, yes otherwise.



---

Deadlock-free  
application layer

+ Deadlock-free  
network layer

+ Deadlock-free  
link layer


?  
=

Deadlock-free system



# NoC Example (2) - Spidergon

---



- Design by STMicroelectronics


# Application layer

---

» Nodes send requests



# Network layer



## Routing logic

$$\text{RelAd} = (\text{dest} - \text{current}) \bmod 4 * N$$

```

if RelAd = 0 then

```

stop

```
elseif 0 < RelAd <= N then
```

go clockwise

```
elseif 3*N <= RelAd <= 4*N then
```

go counter clockwise

**else**


go across

**endif**

# The network has a deadlock !

---

» For instance, we can have a cycle of packets







# Dividing in two networks

---


» Is the system deadlock-free ?



# Dividing in two networks

---


- » Is the system deadlock-free ?
- » Yes ! None of the dependencies in the right upper quarter occur.



# Slaves in 1st quarter only?


---

» Is the system deadlock-free ?




# Slaves in 1st quarter only?

- » Is the system deadlock-free ?
- » No !




# Slaves in 1st quarter only?

- » Is the system deadlock-free ?
- » No !



# Slaves in 1st quarter only?

- » Is the system deadlock-free ?
- » No !





---

Deadlock-free  
application layer

+ Network layer  
with deadlocks

+ Deadlock-free  
link layer

?  
=

Deadlock-free system

# Confusing ...

---

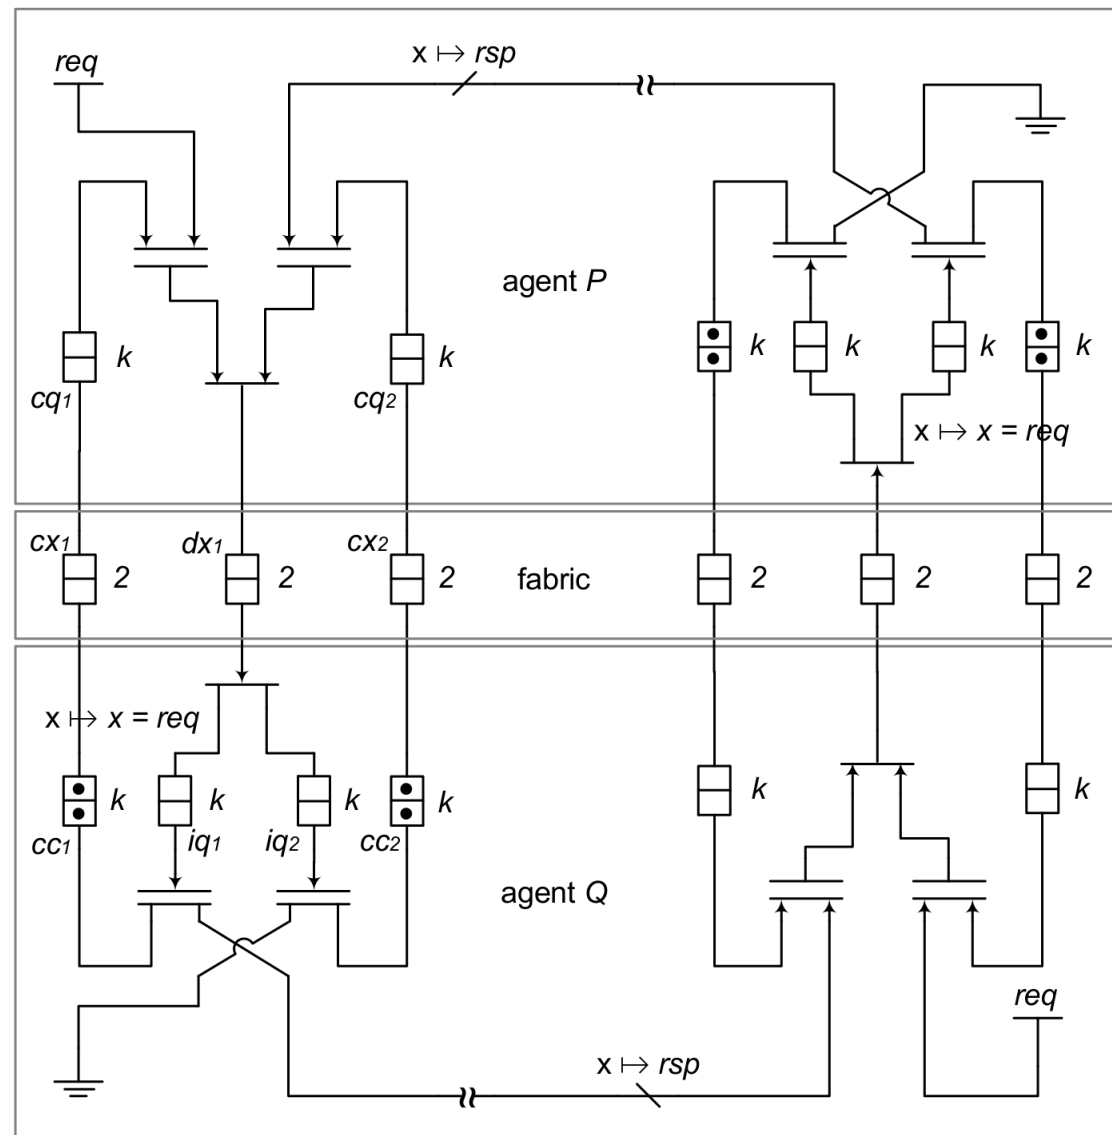
- » We need tools to (quickly) check for deadlocks
  - taking details of **all** three layers into account
  - in **large** systems

# Outline

---

- » Intel's micro-architectural description language
  - xMAS language
  - Capturing high-level structure and message dependencies
  - Extended to “MaDL” at TU/e.
    - Micro-architectural Description Language
- » Deadlock verification for MaDL
  - Definition of deadlocks
  - Labelled dependency graph
  - Feasible logically closed subgraph
- » Conclusion and future work

# Intel's abstraction for networks




- High-level of abstraction
- Exploit high-level structure


Automatic proofs using invariant generation and hardware model-checking

# xMAS

## Executable Micro-Architectural Specification




queue




source




sink




function




fork




join



merge



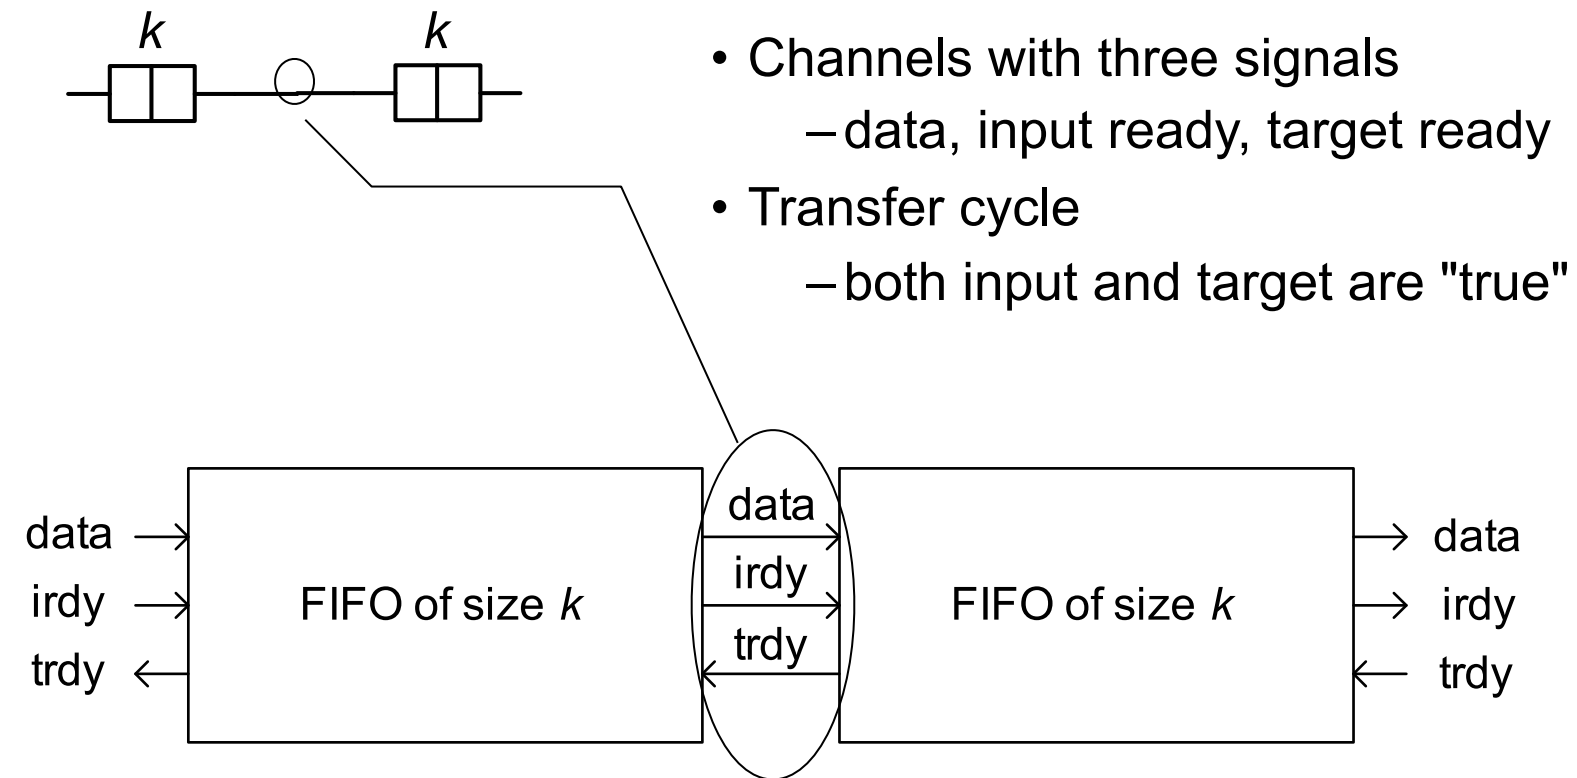

switch



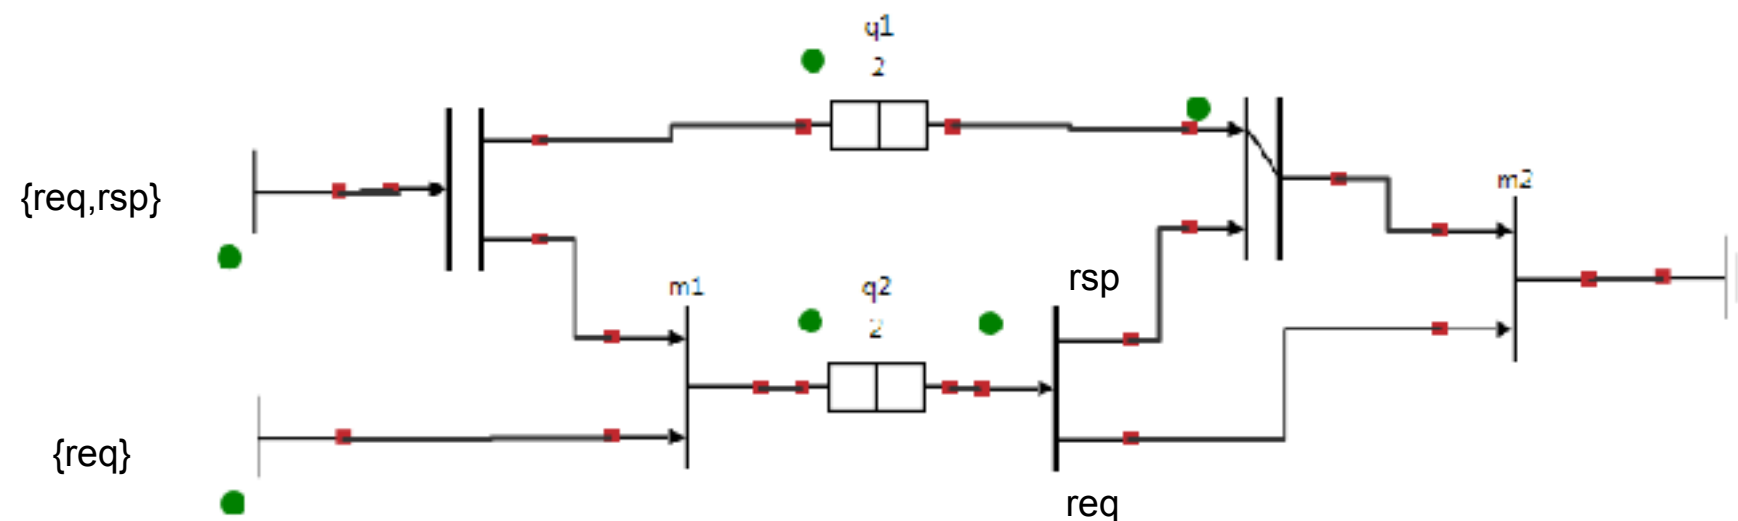
- » Fair sinks and sometimes sources
- » Diagram is formal model
- » Friendly to microarchitects



# MaDL (1)



Notion of a “dead” channel:  $F(c.irdy \ \& \ G \ ! \ c.trdy)$






## MaDL (2)

```


param int QSIZE;
const pkt;
chan src := Source(pkt);
chan one, two := Fork(src);
chan long := Queue(QSIZE, Queue(QSIZE, one));
chan short := Queue(QSIZE, two);
Sink(CtrlJoin(long, short));

```


textual input




IO-FSM



LoadBalancer



Joitch



Multi-Match

additional complex primitives

predicates and functions

```

pred P (p : pkt_t, q : pkt_t) {
    p.id == q.id
};
function F (p : pkt_t) : pkt_t {
    fldA = p.fldB;
    fldB = p.fldA;
};

```

structured recursive data types

```

struct pkt0_t {
    fldA : [2:0];
    fldB : [3:0];
};
union pkt1_t {
    optionA : pktA_t;
    optionB : pktB_t;
};
enum pkt2_t { optionA; optionB; };


const optionA;

```


And more: for-loops, if-then-else, uses, ...

# Composing modules via channels

---



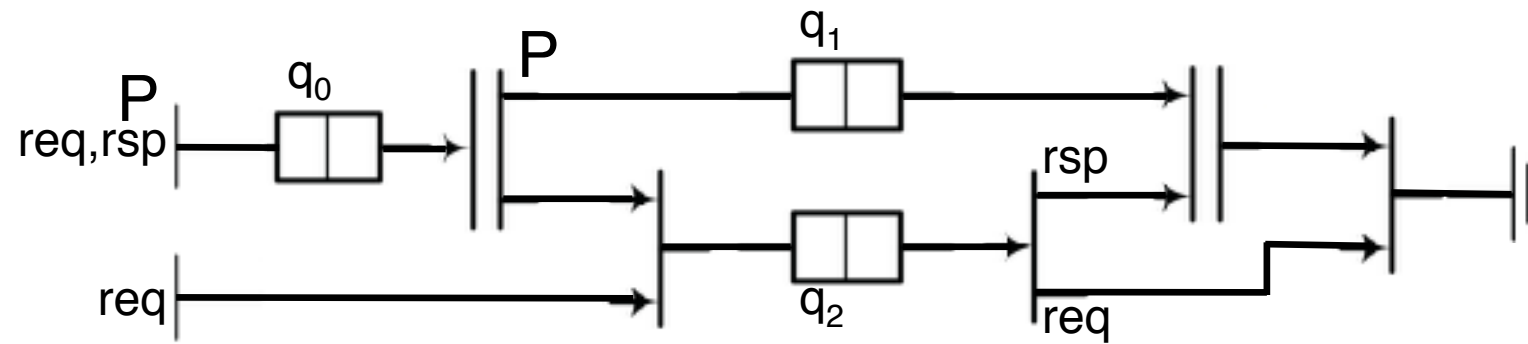
- » Channels with three signals
  - data, input ready, target ready
- » Transfer cycle
  - both input and target are "true"





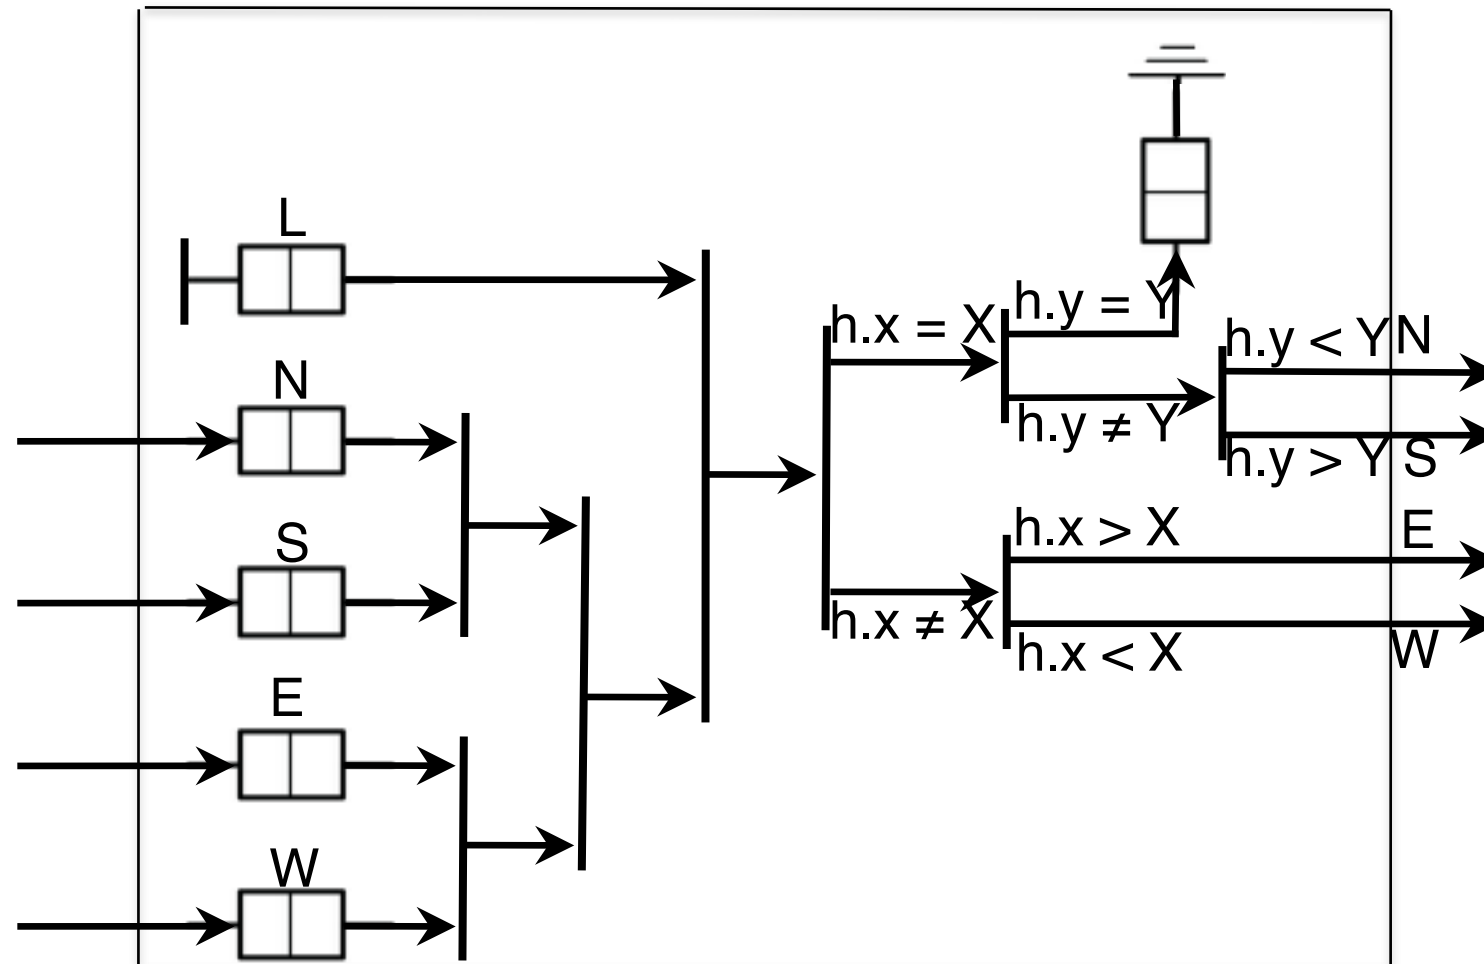
# MaDL example

---

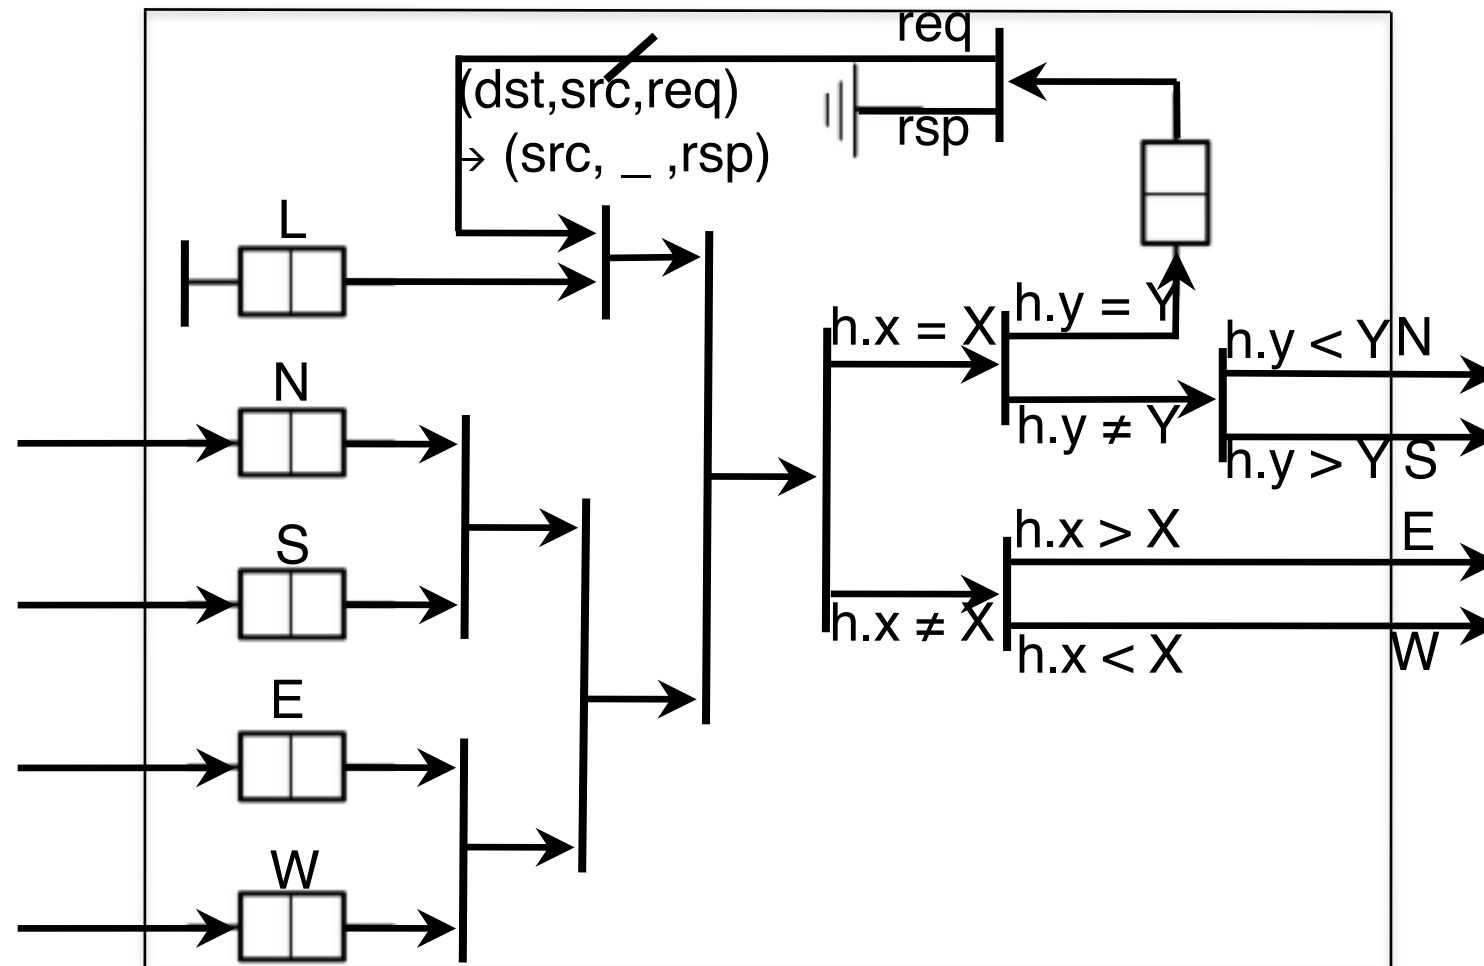


- » Two sources
  - one for requests
  - one for responses

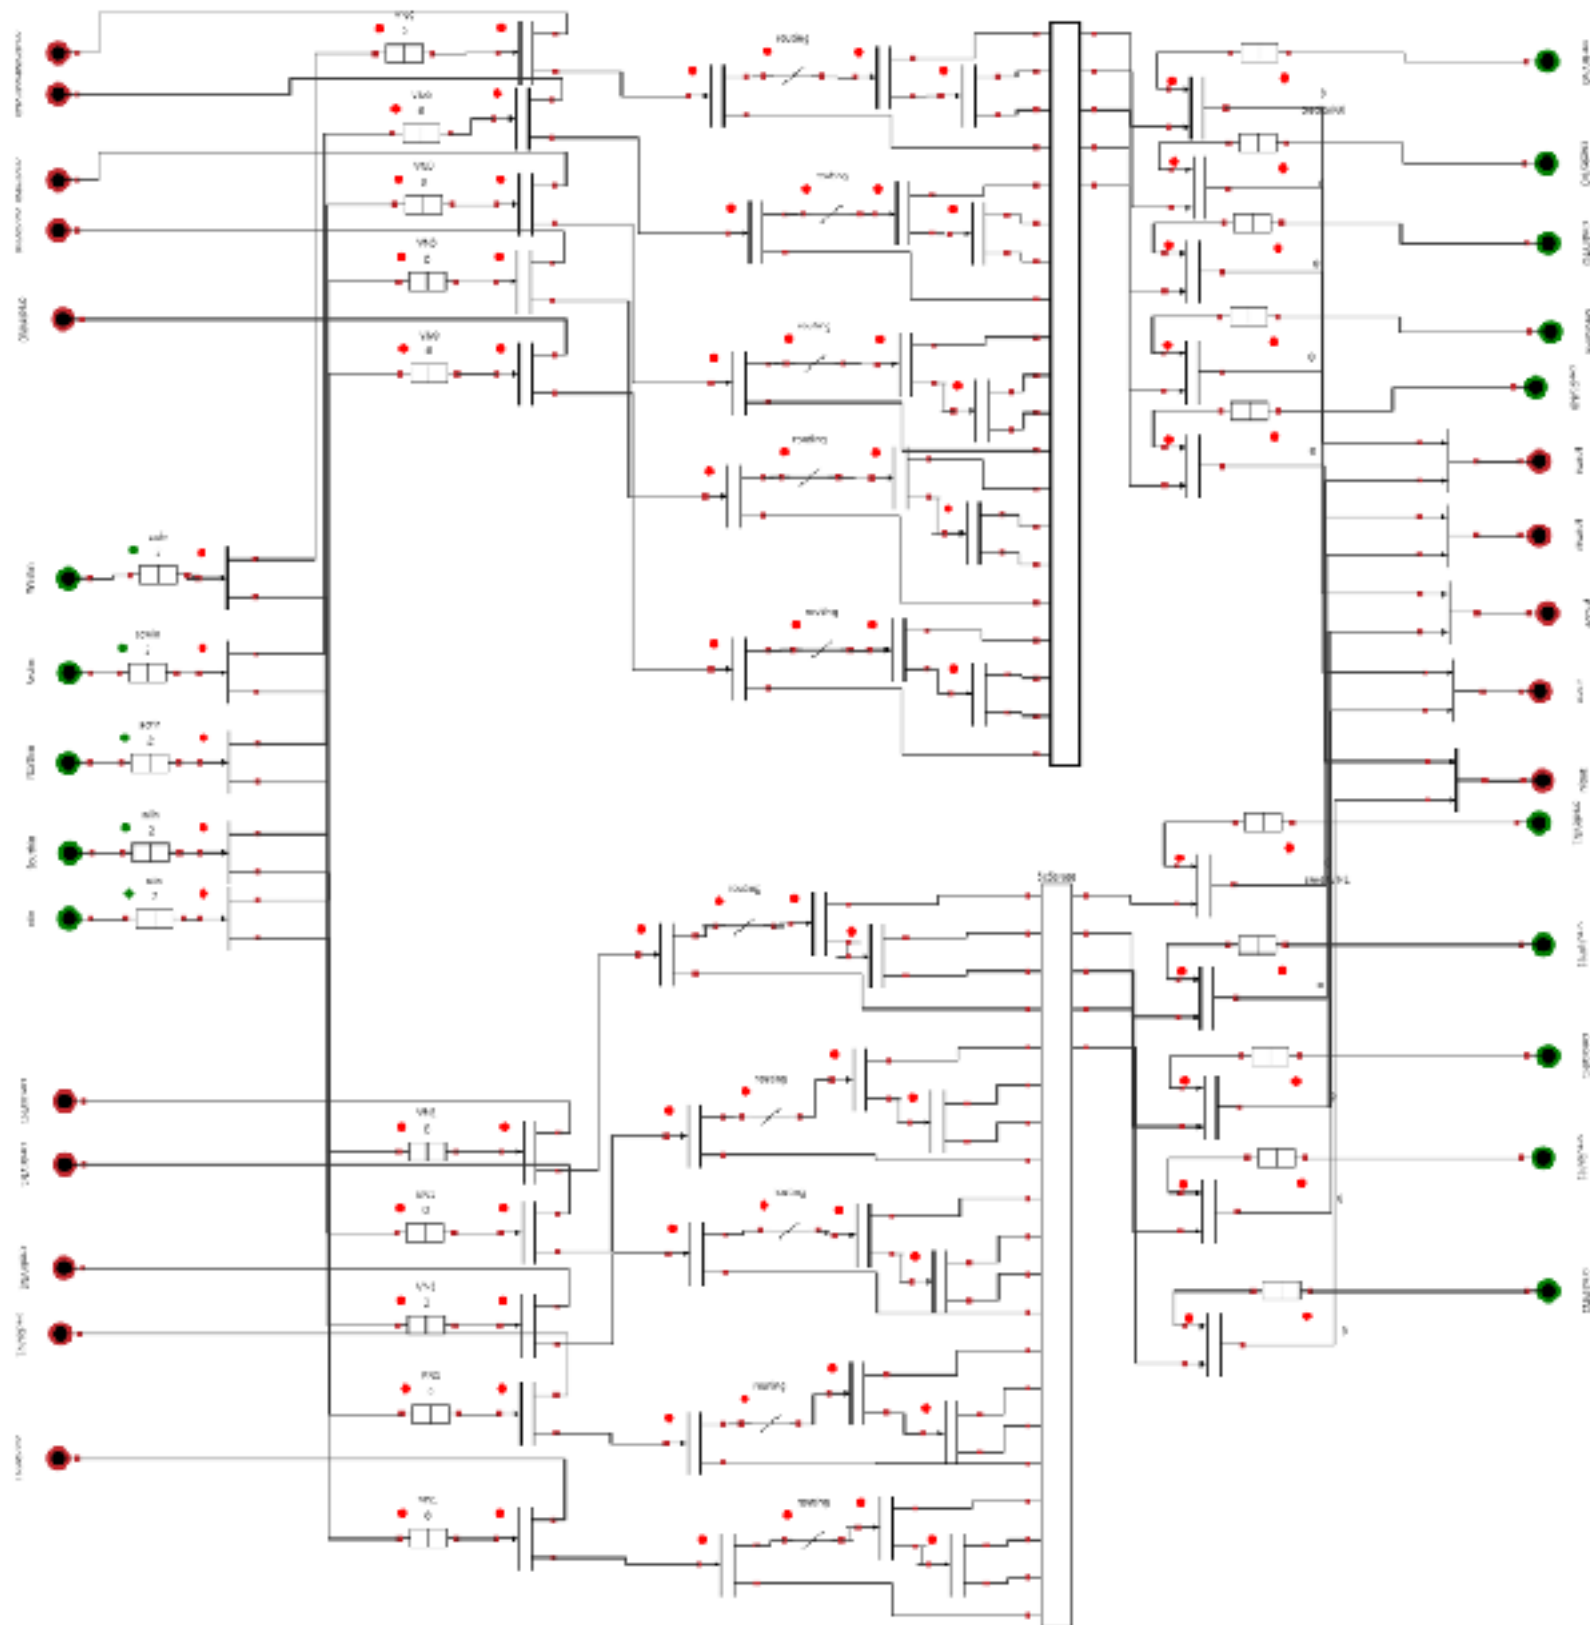
## Processing node for XY routing in a 2D-mesh



## Processing node with requests and responses



## A more complex processing nodes with virtual channels and credits



# Outline

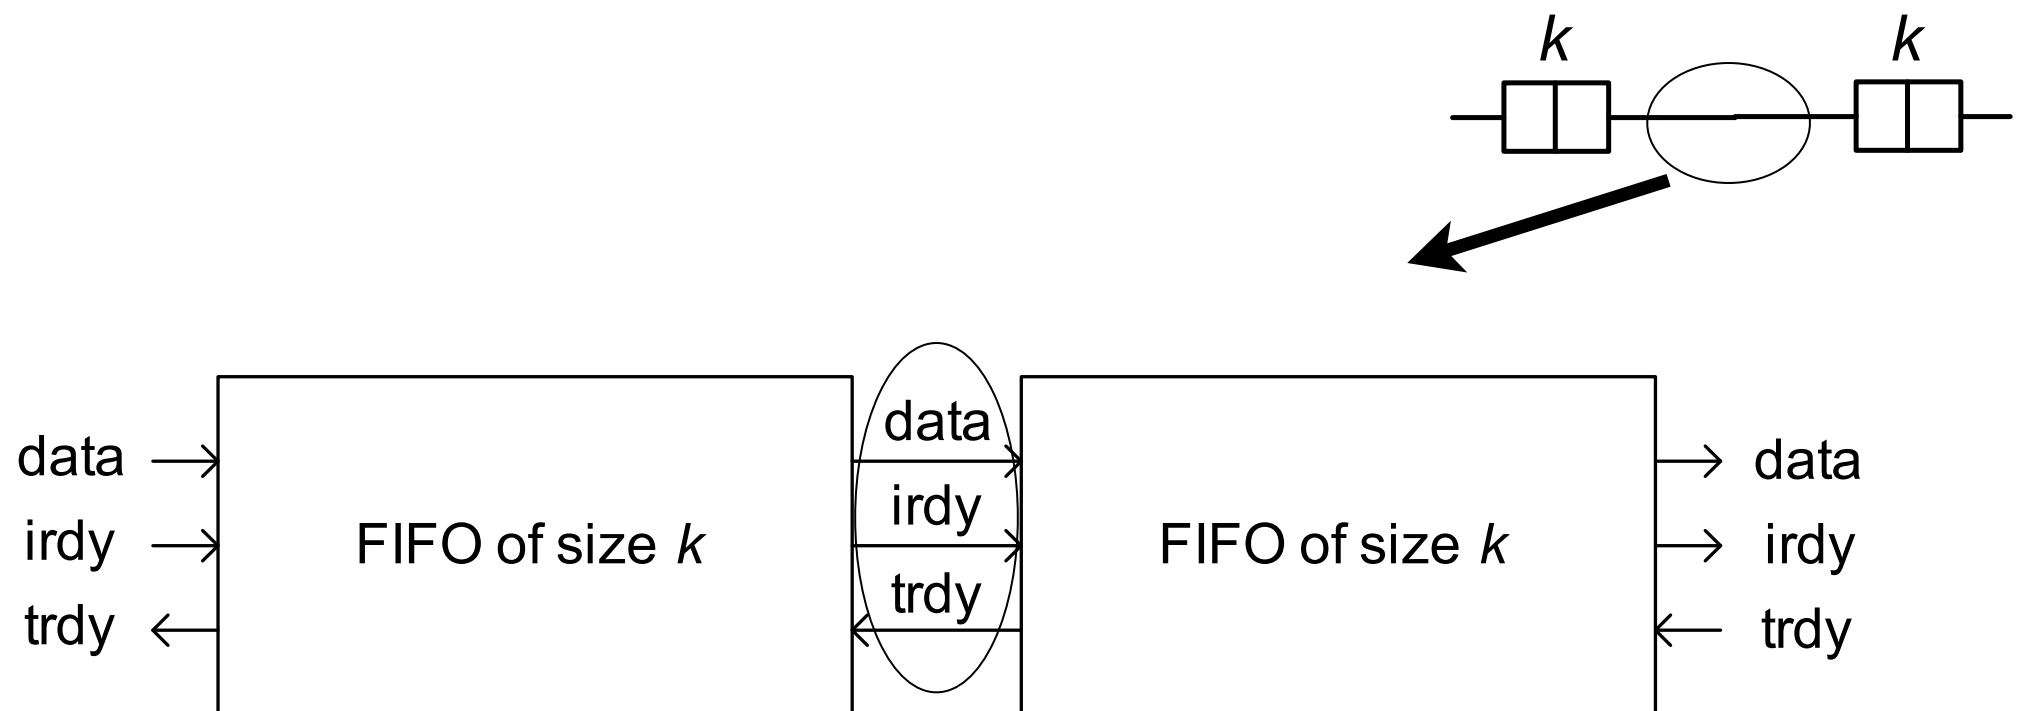
---

- » Intel's micro-architectural description language
  - xMAS language
  - Capturing high-level structure and message dependencies
  - Extended to “MaDL” at TU/e.
    - Micro-architectural Description Language
- » **Deadlock verification for MaDL**
  - **Definition of deadlocks**
  - **Labelled dependency graph**
  - **Feasible logically closed subgraph**
- » Conclusion and future work

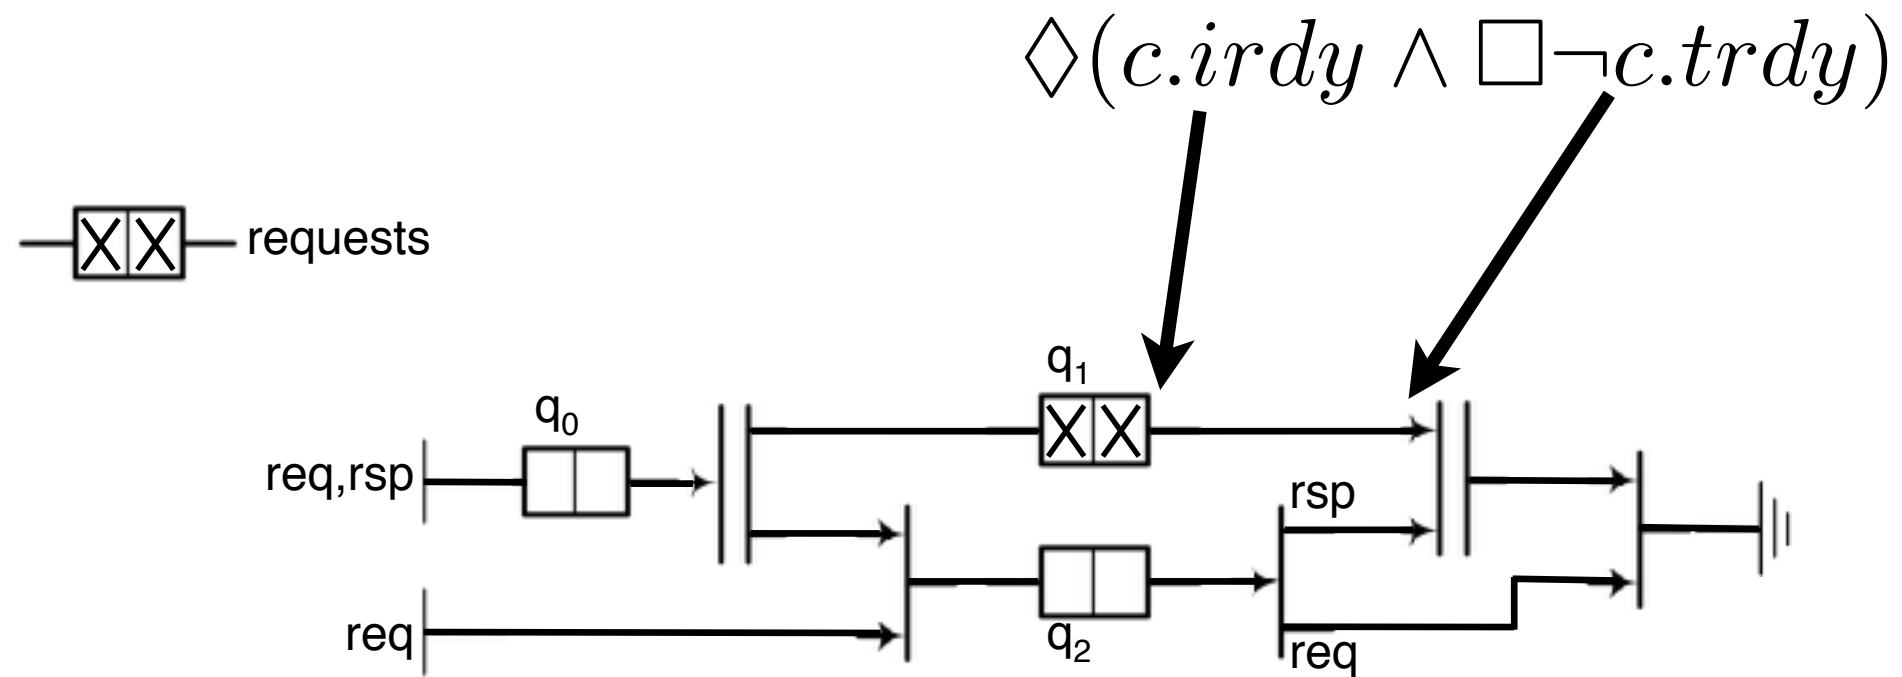
## Formal definition of "deadlock" in MaDL

- Intuition is a "dead" channel
- Formal definition based on Linear Temporal Logic
  - Predicate logic
  - Temporal operators "eventually" ( $\Diamond$ ) and "globally" ( $\Box$ )
- Channel  $c$  is dead iff

$$\Diamond(c.irdy \wedge \Box \neg c.trdy)$$

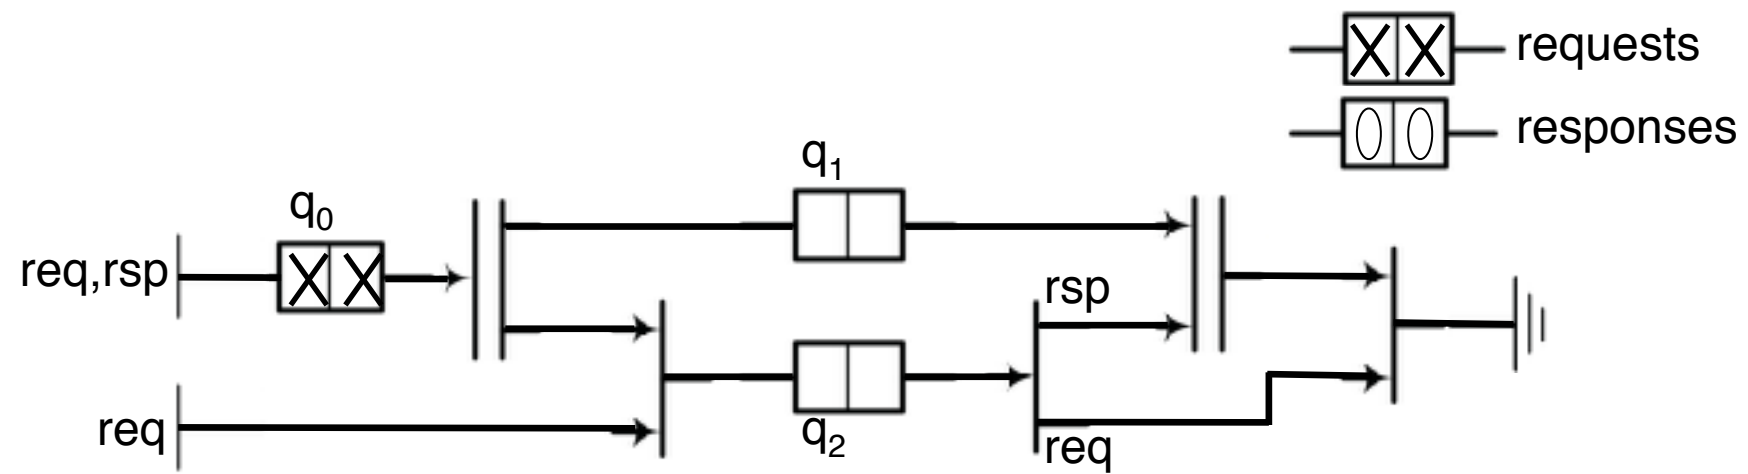


## MaDL example



- Inject two requests in q<sub>0</sub>
- Fork creates two copies
- One pair is sunk

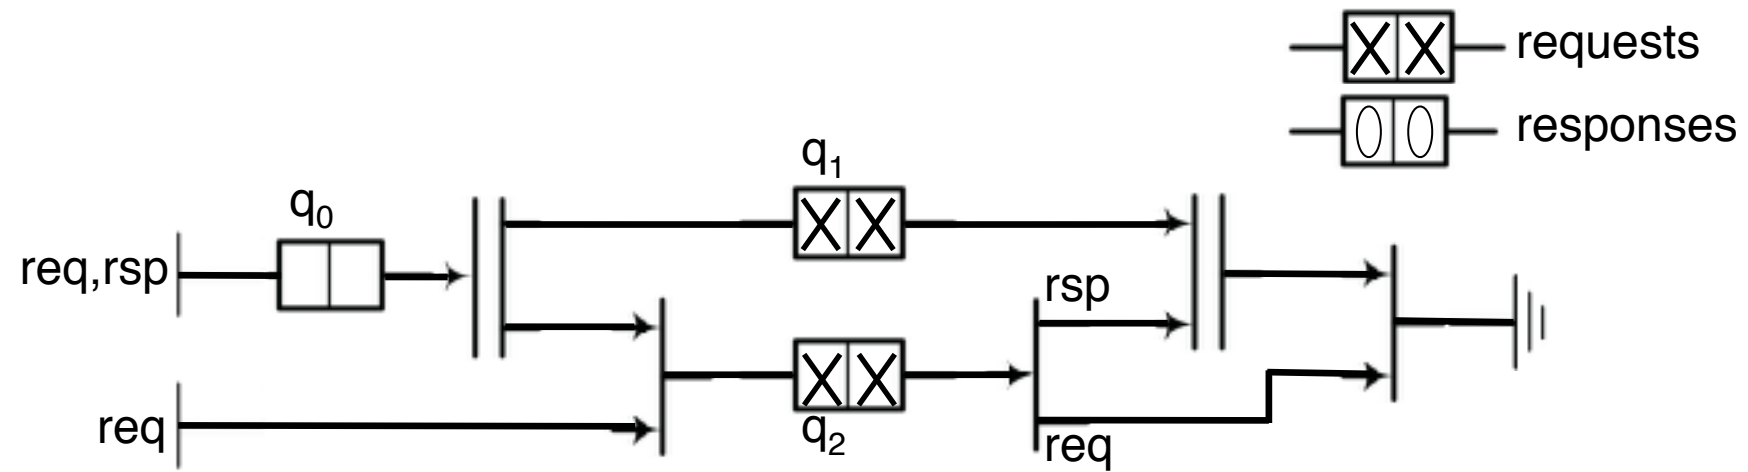
## MaDL example



- Inject two requests in  $q_0$

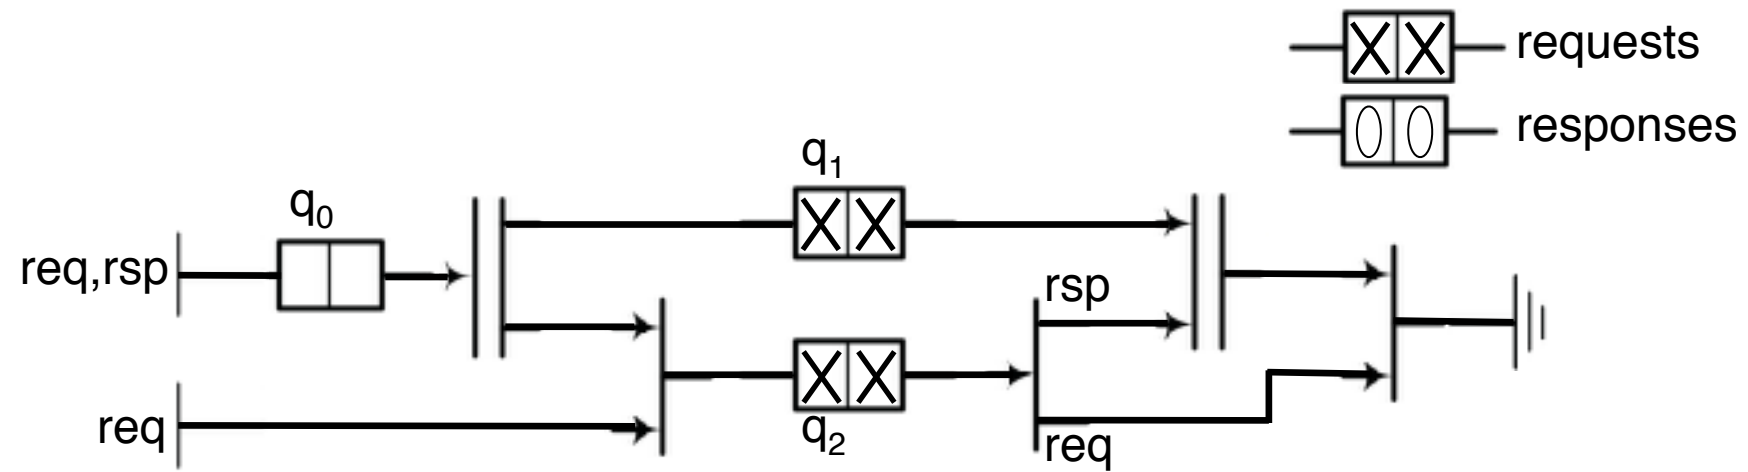


## MaDL example



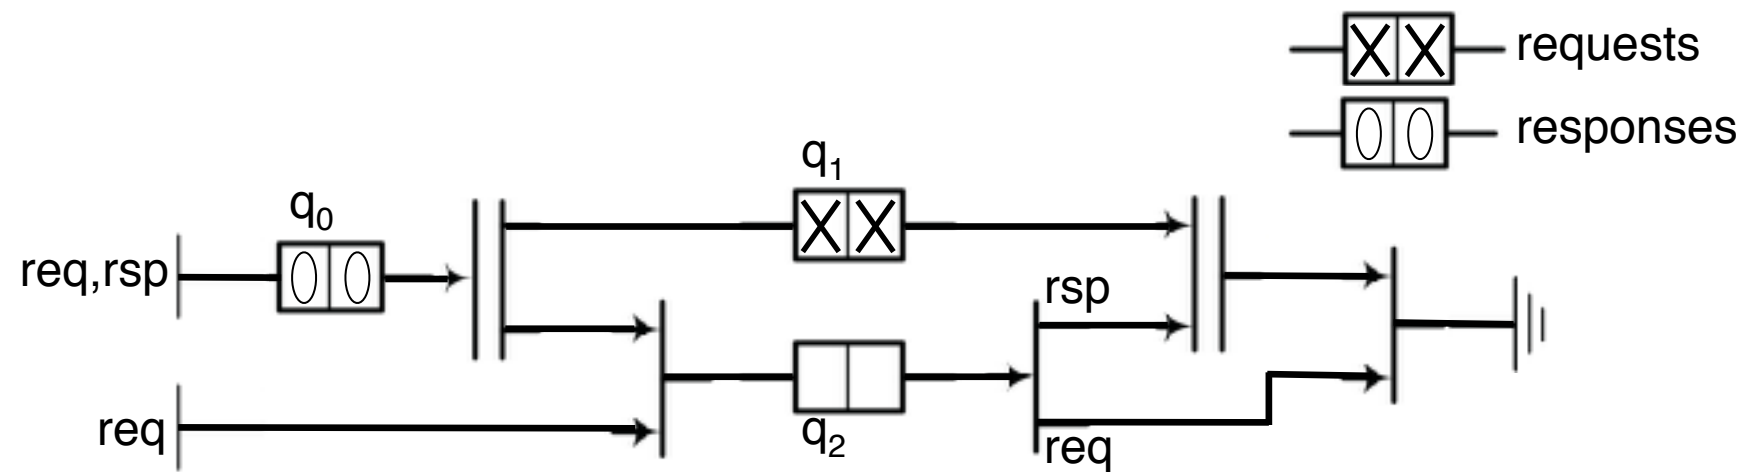
- Inject two requests in  $q_0$
- Fork creates two copies

## MaDL example



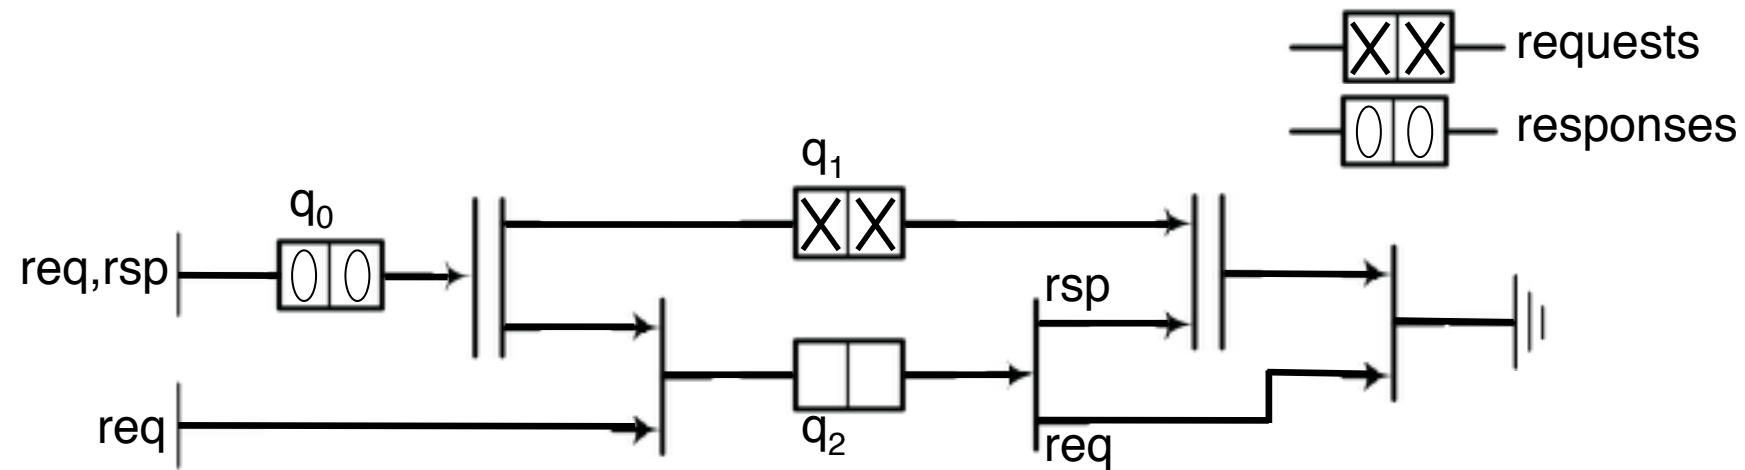
- Inject two requests in q0
- Fork creates two copies
- One pair is sunk

## MaDL example



- Inject two requests in  $q_0$
- Fork creates two copies
- One pair is sunk
- Inject two responses in  $q_0$

## MaDL example



- Inject two requests in  $q_0$
- Fork creates two copies
- One pair is sunk
- Inject two responses in  $q_0$
- $q_2$  is permanently idle for responses,  $q_1$  is permanently blocking

We have a deadlock without a circular wait !

## General approach for deadlock detection in MaDL networks


- Define deadlock equations for all components
  - Equations capture the reason why a component is idle or blocking
- Build a labelled waiting graph for each queue
  - Labels correspond to the equations
  - Graph captures the topology, i.e., the dependencies between the MaDL components
- Search for a feasible logically closed subgraph
  - Corresponds to a deadlock situation
  - Feasibility checked using Linear Programming
- This approach may output unreachable deadlocks
  - A first step generates invariants to rule out false deadlocks
  - Invariants are rather weak and simple - false deadlocks are in theory still possible

## General approach for deadlock detection in xMAS networks

- Define deadlock equations for all components
  - Equations capture the reason why a component is idle or blocking
- Build a labelled waiting graph for each queue
  - Labels correspond to the equations
  - Graph captures the topology, i.e., the dependencies between the MaDL components
- Search for a feasible logically closed subgraph
  - Corresponds to a deadlock situation
  - Feasibility checked using Linear Programming
- This approach may output unreachable deadlocks
  - A first step generates invariants to rule out false deadlocks
  - Invariants are rather weak and simple - false deadlocks are in theory still possible

## Deadlock equations for a channel


- Depends on the **target** component connected to the channel
- We look at the input port of the target component



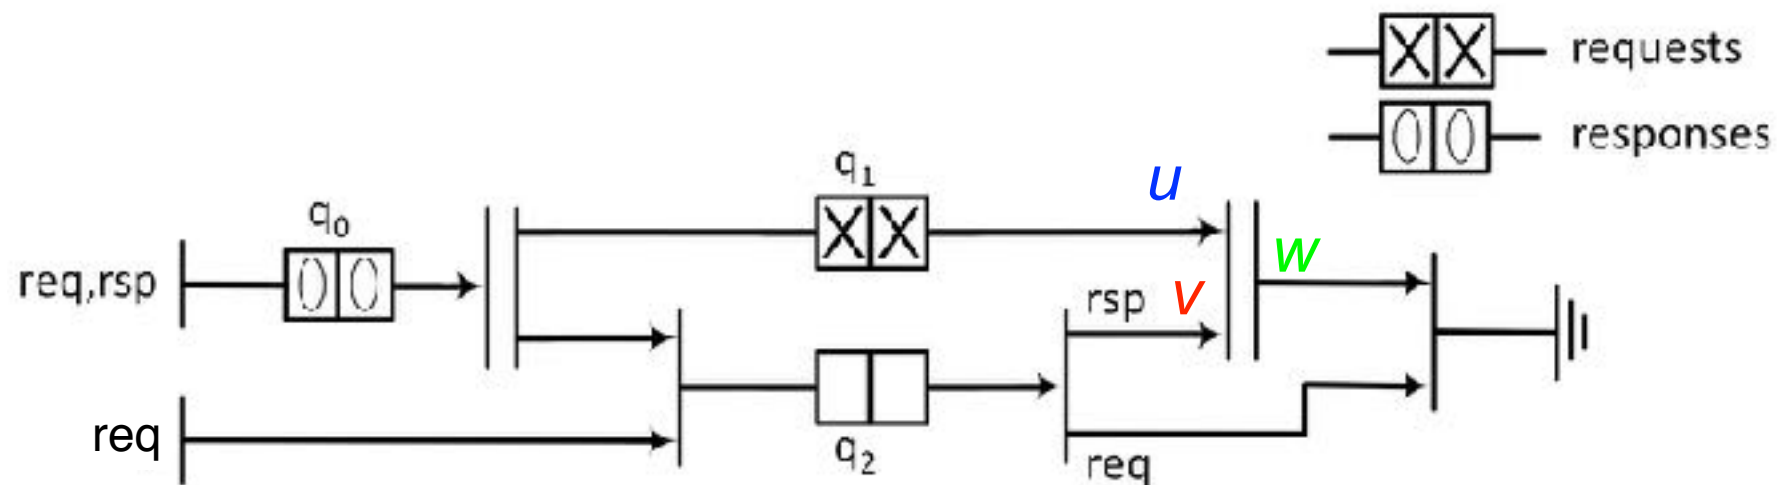
## Deadlock equations for a join

- 2 cases
  - output is blocked
  - the other input is idle

We need to know when a channel is idle !




- $\text{Block}(u) = \text{Idle}(v) + \text{Block}(w)$





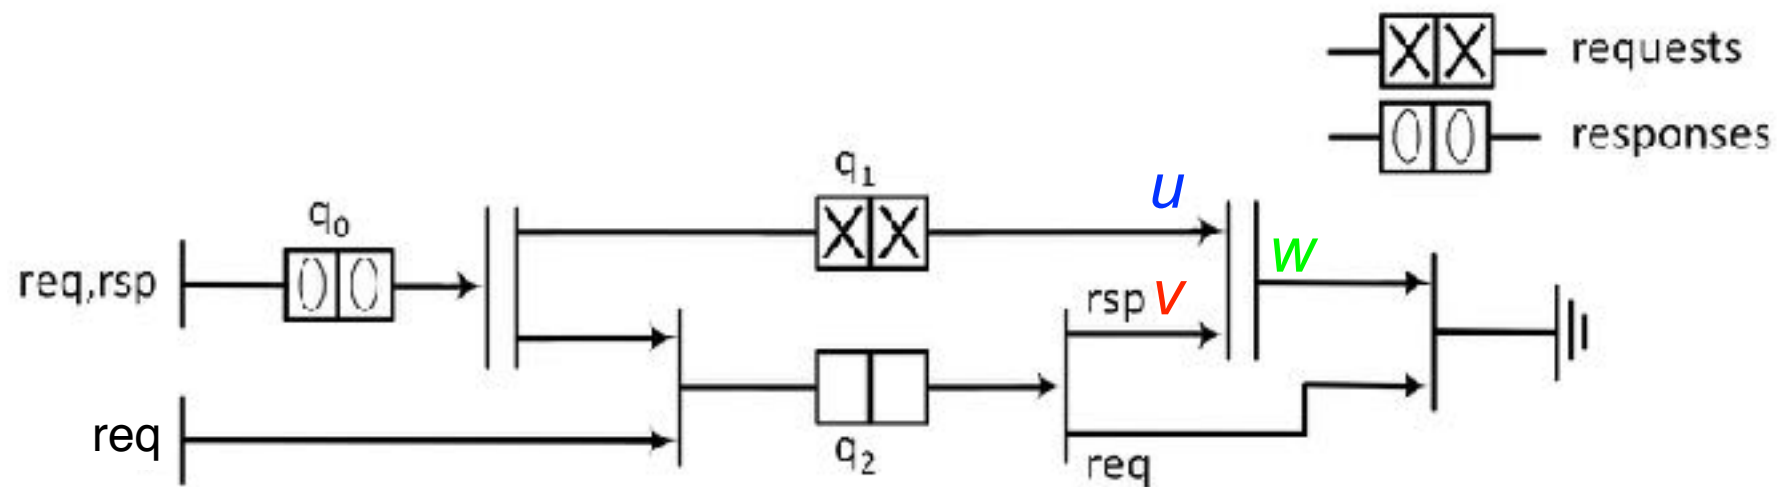

## Idle equations for a channel

- Depends on the **initiator** component connected to the channel
- We are looking at the input port of the initiator



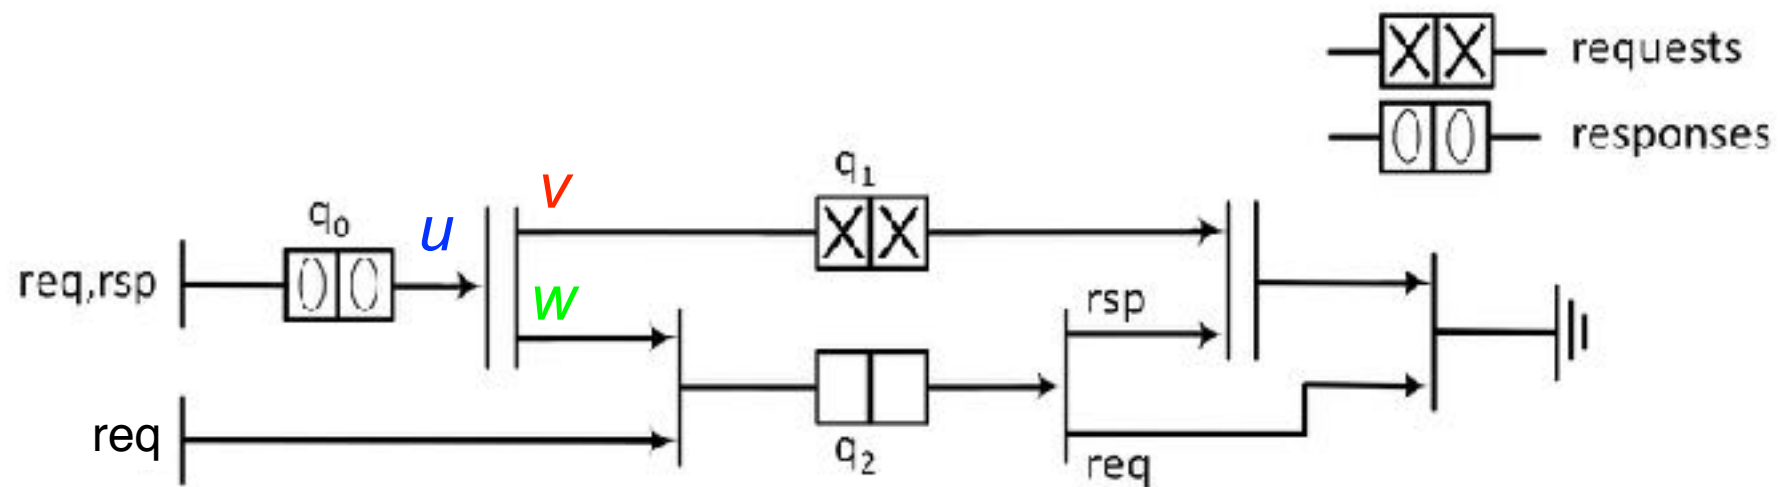

## Idle equations for a join

- A join is idle if one of the input channels is idle
- $\text{Idle}(w) = \text{Idle}(u) + \text{Idle}(v)$



## Idle equations for a fork

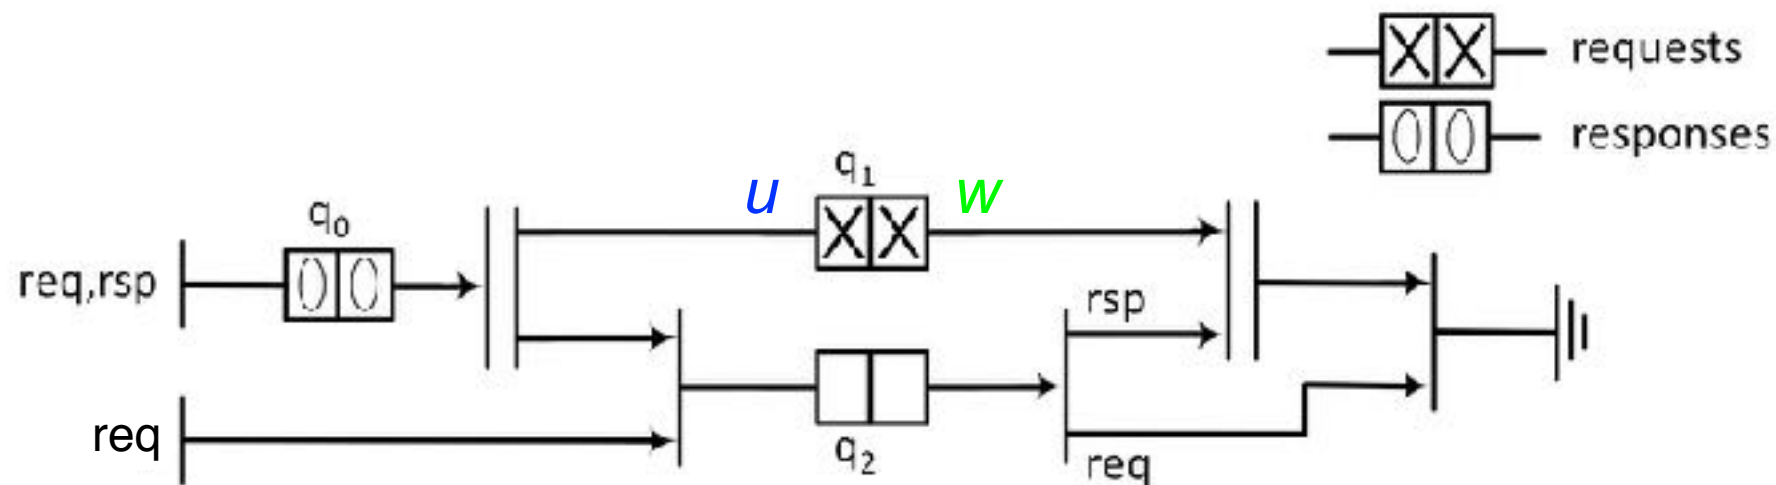

- A fork output is idle if the input is idle or the other output is blocked
- $\text{Idle}(w) = \text{Idle}(u) + \text{Block}(v)$



## Idle equations for a queue

- A queue is idle if it is empty and its input channel is idle
- This is for one message type which might be blocked by another type

- $\text{Idle}(w) = \text{Empty}(q) \cdot \text{Idle}(u) + \text{Block}(w')$   
 –where  $w'$  is a message with a type different from  $w$




## Our quest for "dead" queues

- Definition of a deadlock
  - $\neg F (u.irdy \Rightarrow G \sim u.trdy)$
- We look for a "dead" queue
  - with a message in it ( $u.irdy$ )
  - output blocked ( $G \sim u.trdy$ )
- Over approximation
  - configuration not always reachable
  - we may output false deadlocks



# MaDL Verification: Overview



Generate Invariants

Generate SMT Problem

SMT  
Solver (Z3)

(unsat)



No deadlock!

(sat)



Generate SMV Model

nuXmv



Deadlock!


Basic idea:

- Express deadlock in SMT
- Over-approximate deadlocks
- Use invariants to rule-out false deadlocks
- Reachability analysis of found deadlocks

## General approach for deadlock detection in MaDL networks


- Define deadlock equations for all components
  - Equations capture the reason why a component is idle or blocking
- Build a labelled waiting graph for each queue
  - Labels correspond to the equations
  - Graph captures the topology, i.e., the dependencies between the MaDL components
- Search for a feasible logically closed subgraph
  - Corresponds to a deadlock situation
  - Feasibility checked using Linear Programming
- This approach may output unreachable deadlocks
  - A first step generates invariants to rule out false deadlocks
  - Invariants are rather weak and simple - false deadlocks are in theory still possible

## Step 1 / simulation - req







## Step 1 / simulation - **rsp**



## General approach for deadlock detection in MaDL networks

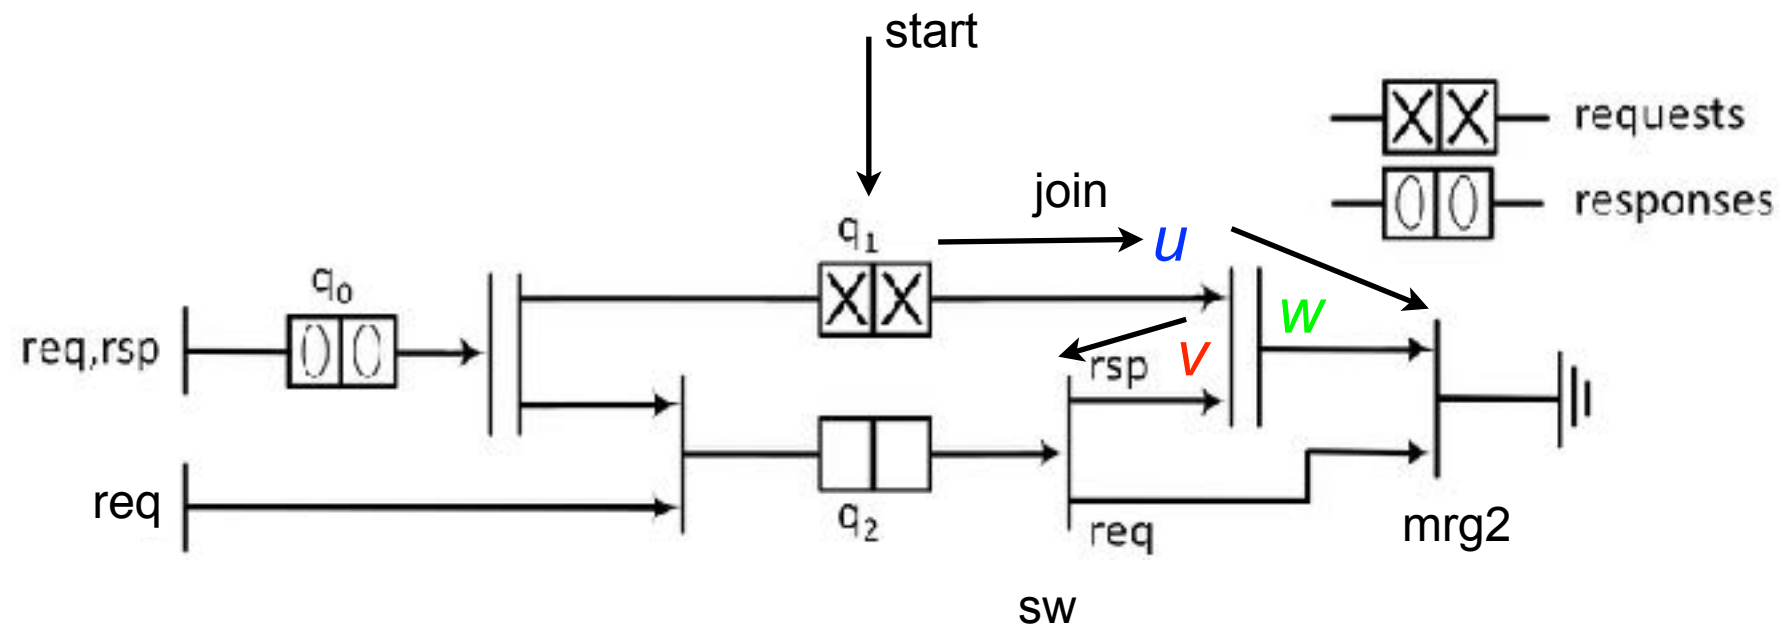
- Define deadlock equations for all components
  - Equations capture the reason why a component is idle or blocking
- Build a labelled waiting graph for each queue
  - Labels correspond to the equations
  - Graph captures the topology, i.e., the dependencies between the MaDL components
- Search for a feasible logically closed subgraph
  - Corresponds to a deadlock situation
  - Feasibility checked using Linear Programming
- This approach may output unreachable deadlocks
  - A first step generates invariants to rule out false deadlocks
  - Invariants are rather weak and simple - false deadlocks are in theory still possible

## Step 2 / labelled dependency graph (1)



start with a message in q1 and visit the join


## Step 2 / labelled dependency graph (2)



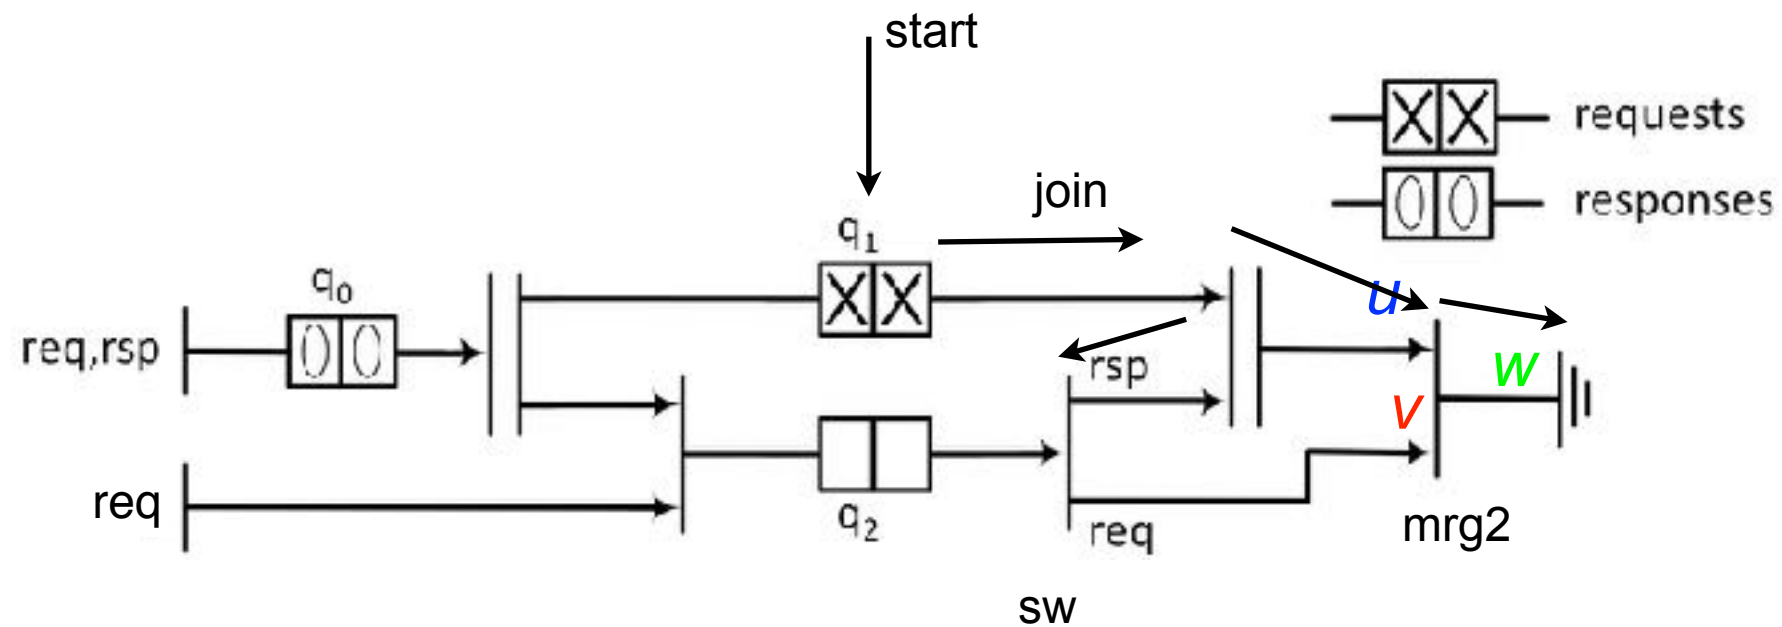
$$\mathbf{Block}(u) = \mathbf{Idle}(v) + \mathbf{Block}(w)$$

analyse the join according to its deadlock equation

we go forward to the merge and backward to the switch




## Step 2 / labelled dependency graph (2)



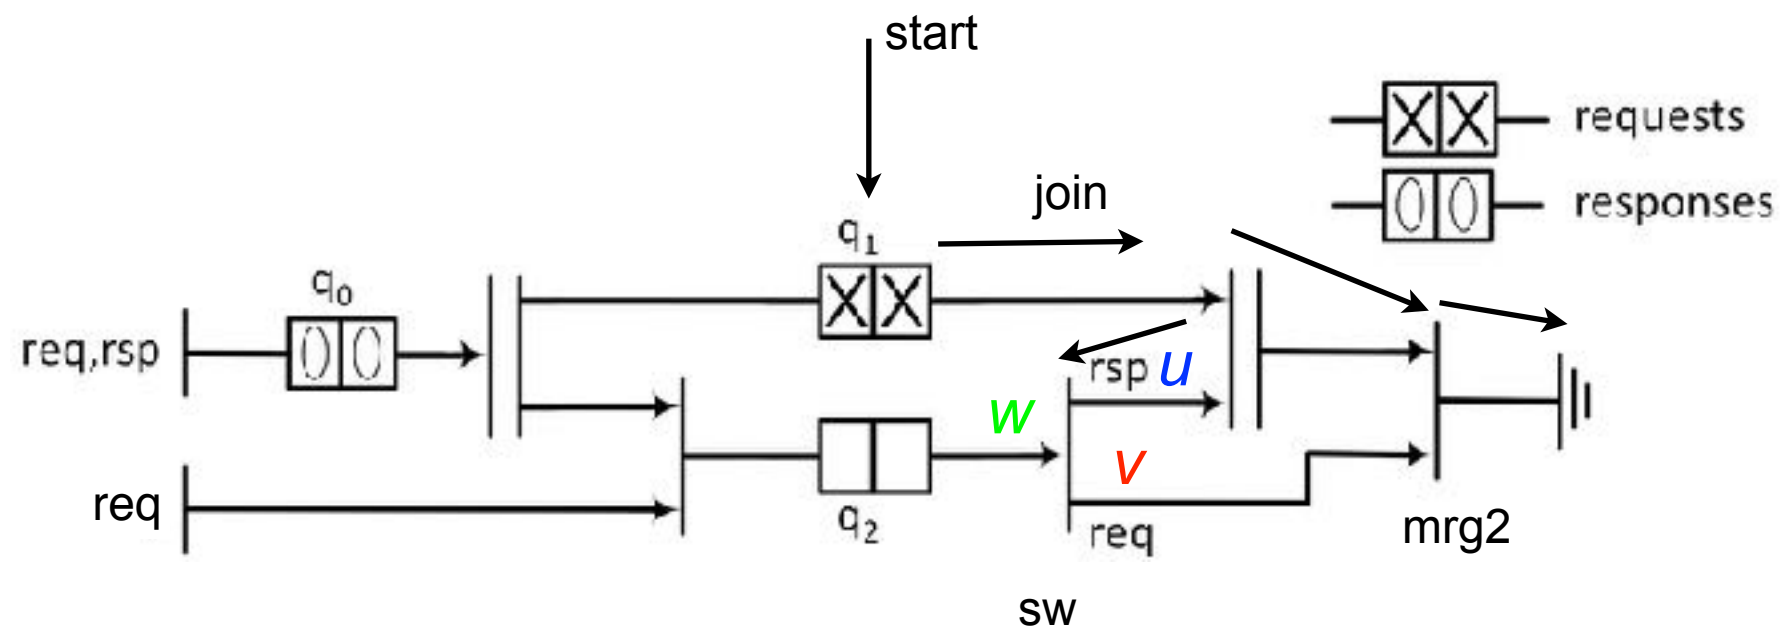
**Block(u) = Block(w)**

forwards to the switch - then the sink can never be blocked

we assume fair sinks




## Step 2 / labelled dependency graph (2)

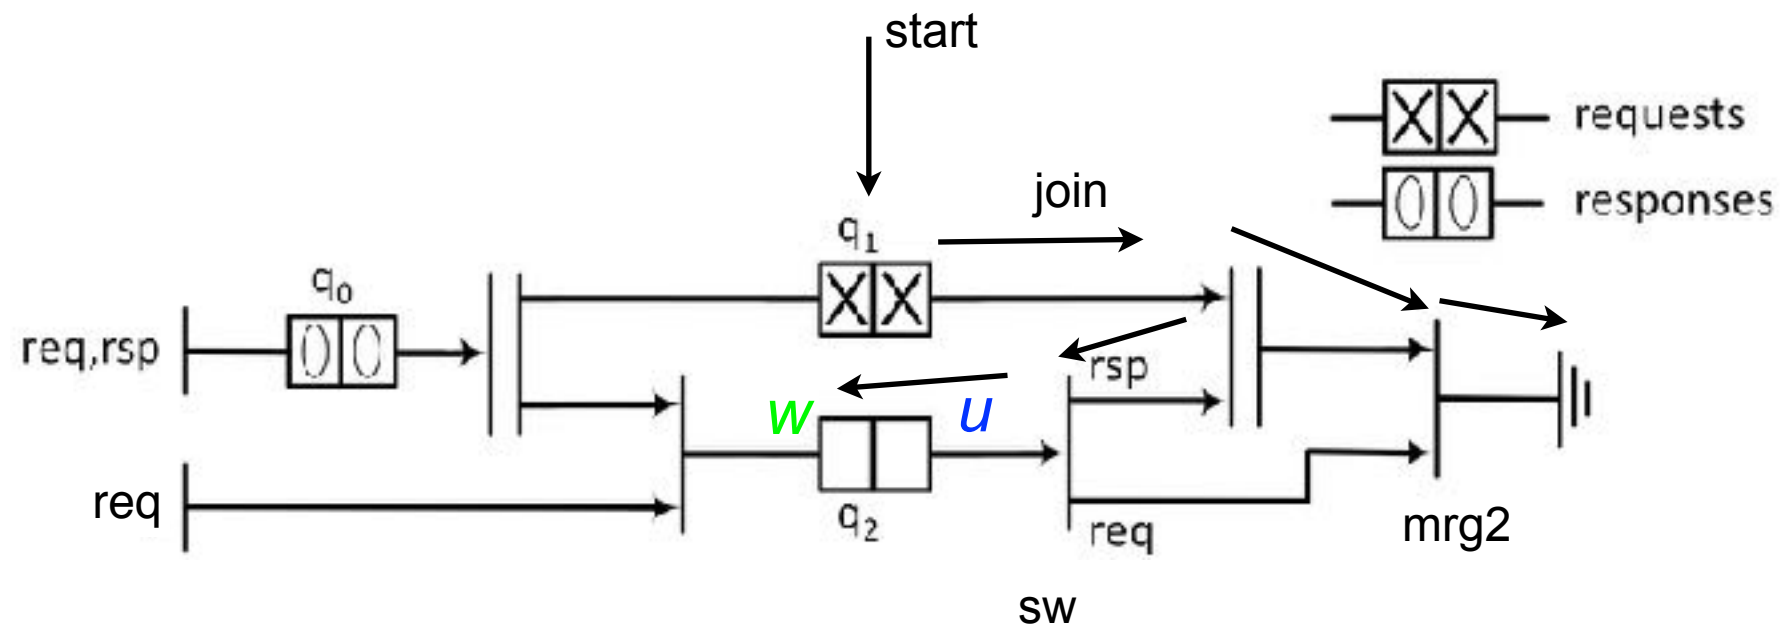


$$\text{Idle}(u) = \text{Idle}(w)$$

backwards to the switch




## Step 2 / labelled dependency graph (2)



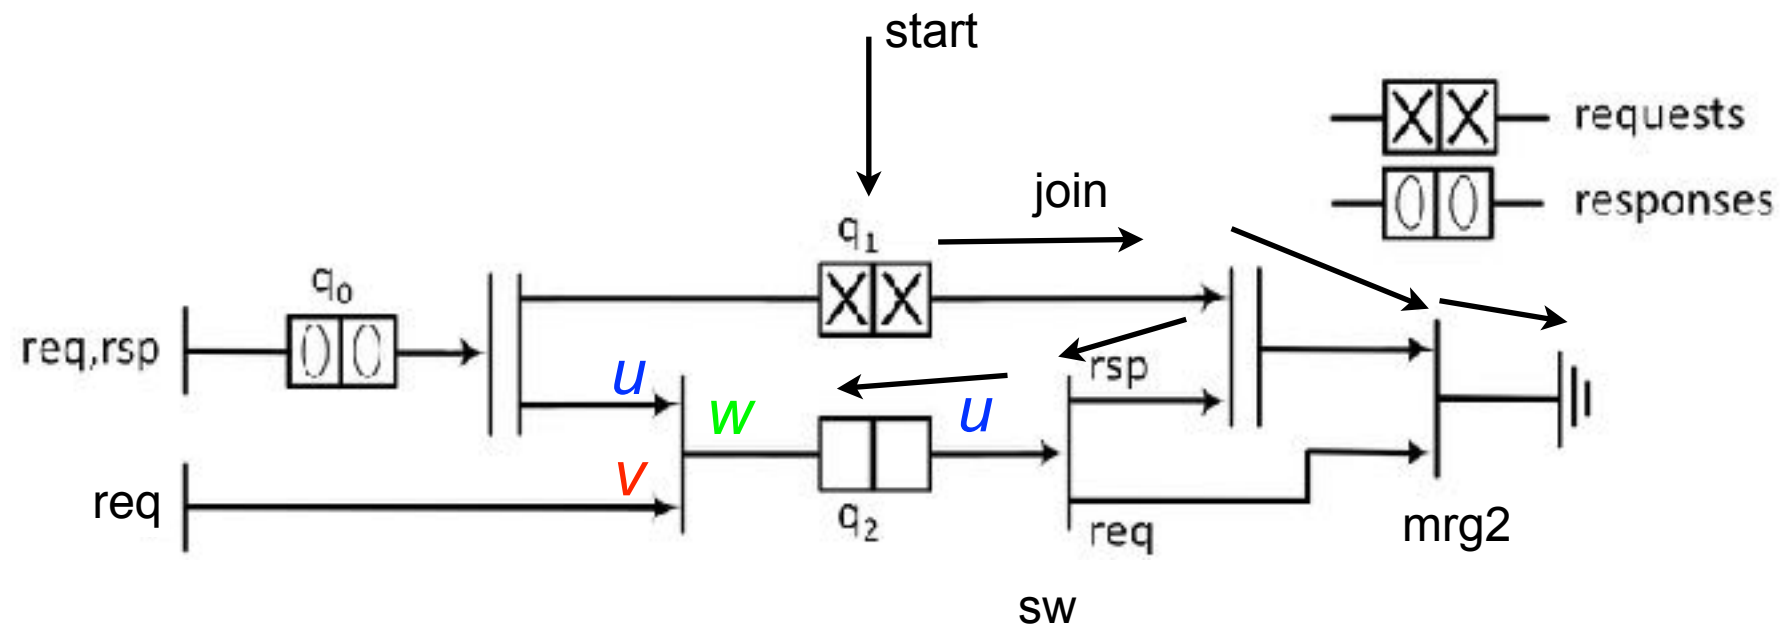
$$\text{Idle}(u) = \text{Idle}(w) \cdot \text{Empty}(q_2)$$

backwards to the queue

note that we forgot the **Block**( $w'$ ) case




## Step 2 / labelled dependency graph (2)



$$\text{Idle}(w) = \text{Idle}(u) \cdot \text{Idle}(v)$$

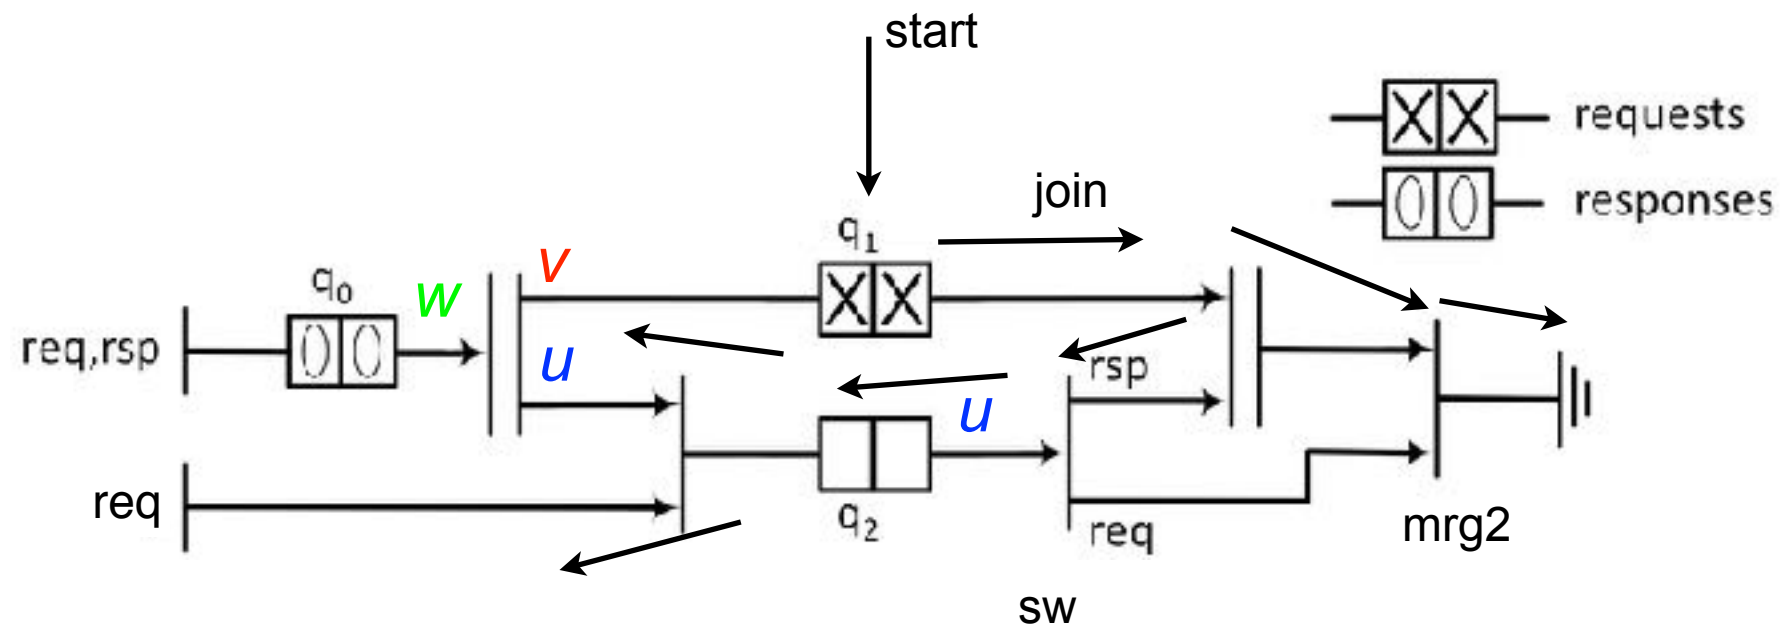
backwards to the merge and branch

note branching is bad for us






## Step 2 / labelled dependency graph (2)

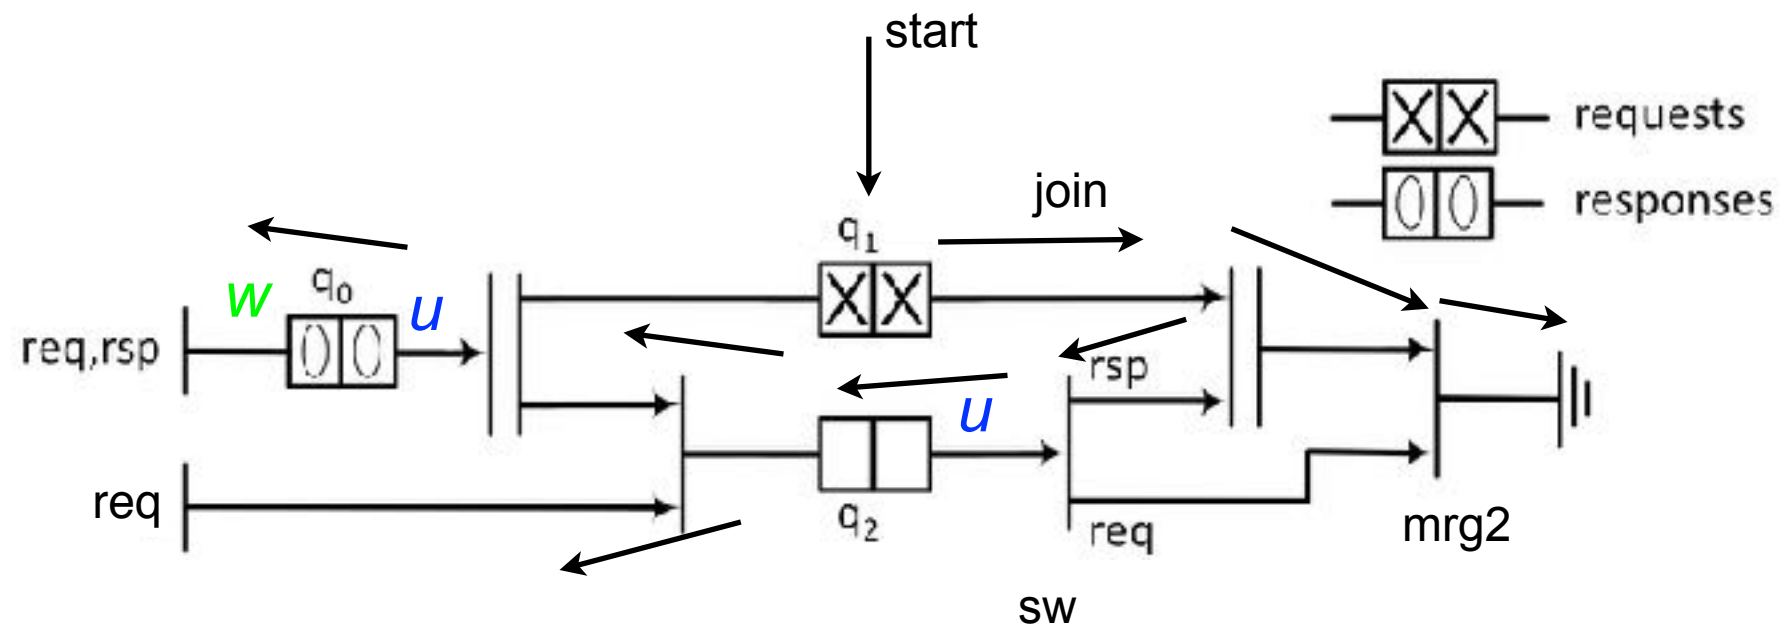


$$\text{Idle}(u) = \text{Block}(v) + \text{Idle}(w)$$

backwards to the merge and branch  
to the source - idle if no type produced  
to the fork




## Step 2 / labelled dependency graph (2)

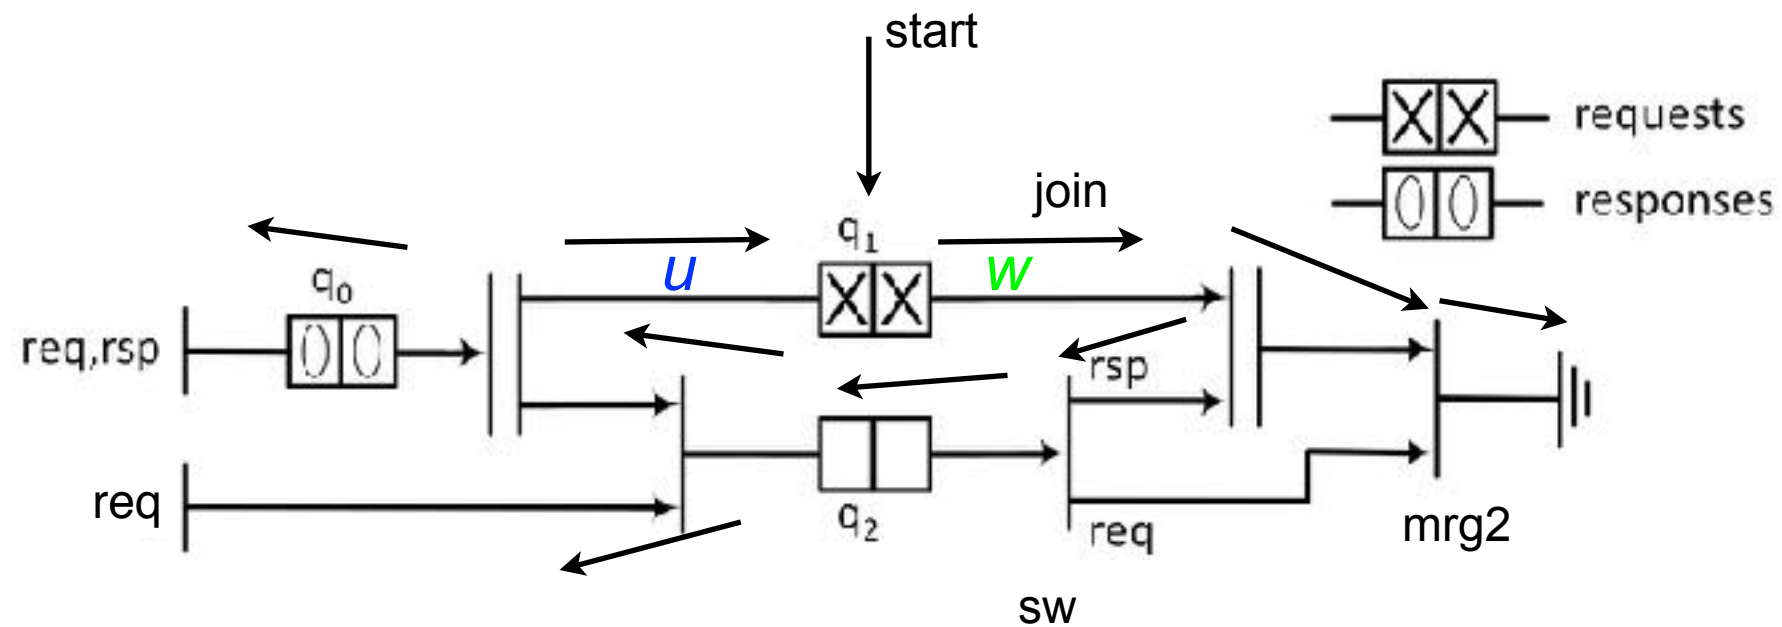


$$\text{Idle}(u) = \text{Idle}(w) \cdot \text{Empty}(q_0)$$

backwards to  $q_0$  and the source




## Step 2 / labelled dependency graph (2)



$$\mathbf{Block}(u) = \mathbf{Block}(w) \cdot \mathbf{Full}(q_1)$$


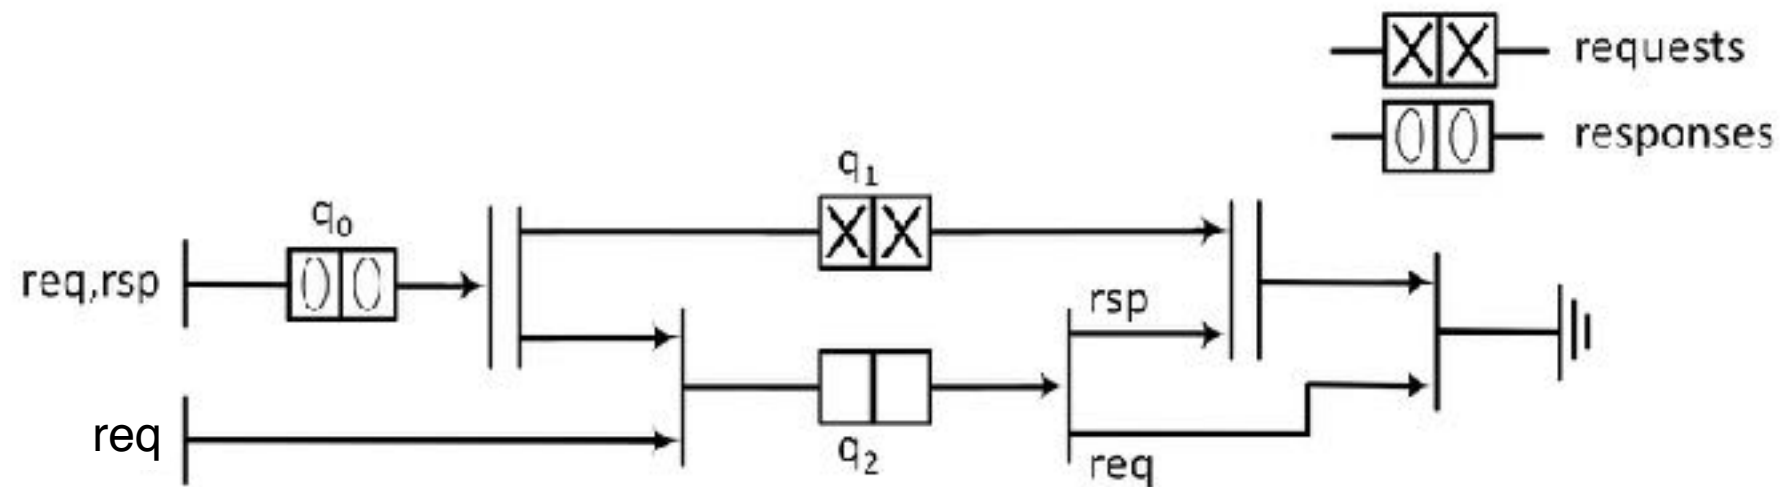
forwards back to  $q_1$  and stop expansion




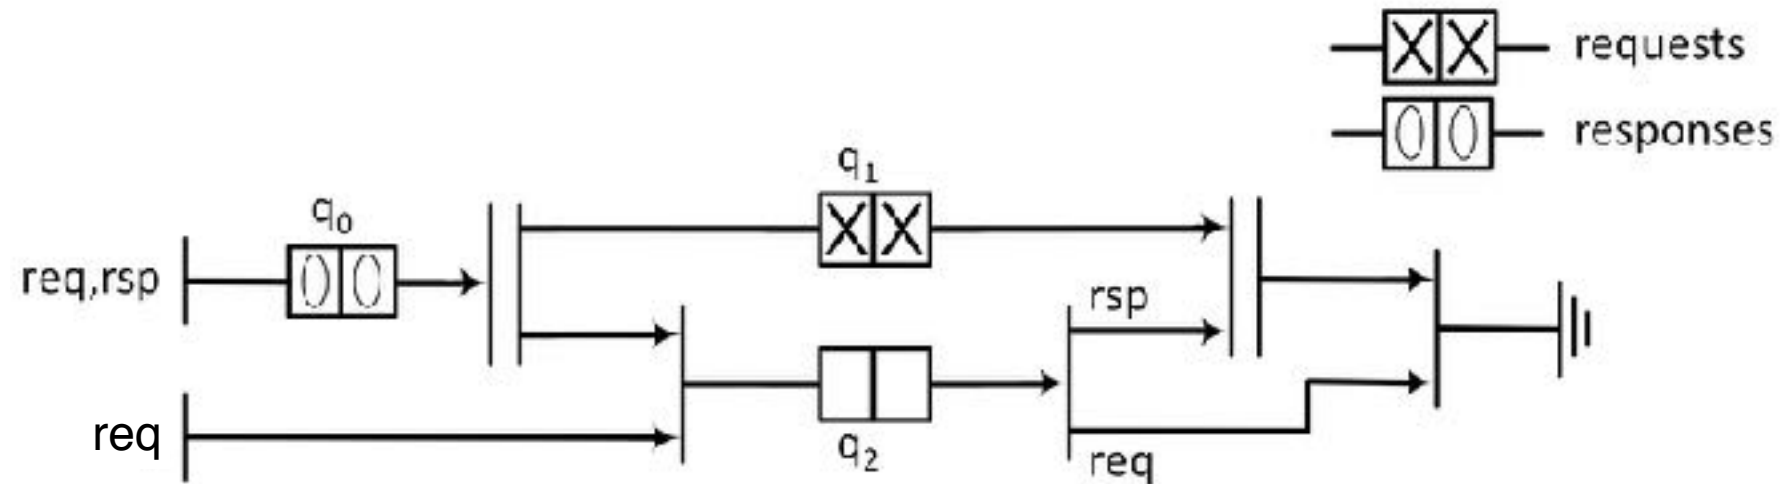
## General approach for deadlock detection in MaDL networks

- Define deadlock equations for all components
  - Equations capture the reason why a component is idle or blocking
- Build a labelled waiting graph for each queue
  - Labels correspond to the equations
  - Graph captures the topology, i.e., the dependencies between the MaDL components
- Search for a feasible logically closed subgraph
  - Corresponds to a deadlock situation
  - Feasibility checked using Linear Programming
- This approach may output unreachable deadlocks
  - A first step generates invariants to rule out false deadlocks
  - Invariants are rather weak and simple - false deadlocks are in theory still possible

## Step 2 / logically closed subgraph 1



## Step 2 / logically closed subgraph 2






# Invariant Generation (1)

Flow invariants:

- Automatically generated
- Linear equations over the # of packets in channels
  - Gaussian elimination
  - Equalities of # of packet between queues
  - $\text{NumX} = \# \text{In} - \# \text{Out}$



Linear equations:

- $\# \text{in} = \# \text{A\_in} = \# \text{B\_in}$
- $\# \text{A} = \# \text{A\_in} - \# \text{A\_out}$
- $\# \text{B} = \# \text{B\_in} - \# \text{B\_out}$
- $\# \text{A\_out} = \# \text{B\_out} = \# \text{out}$

After Gaussian elimination:

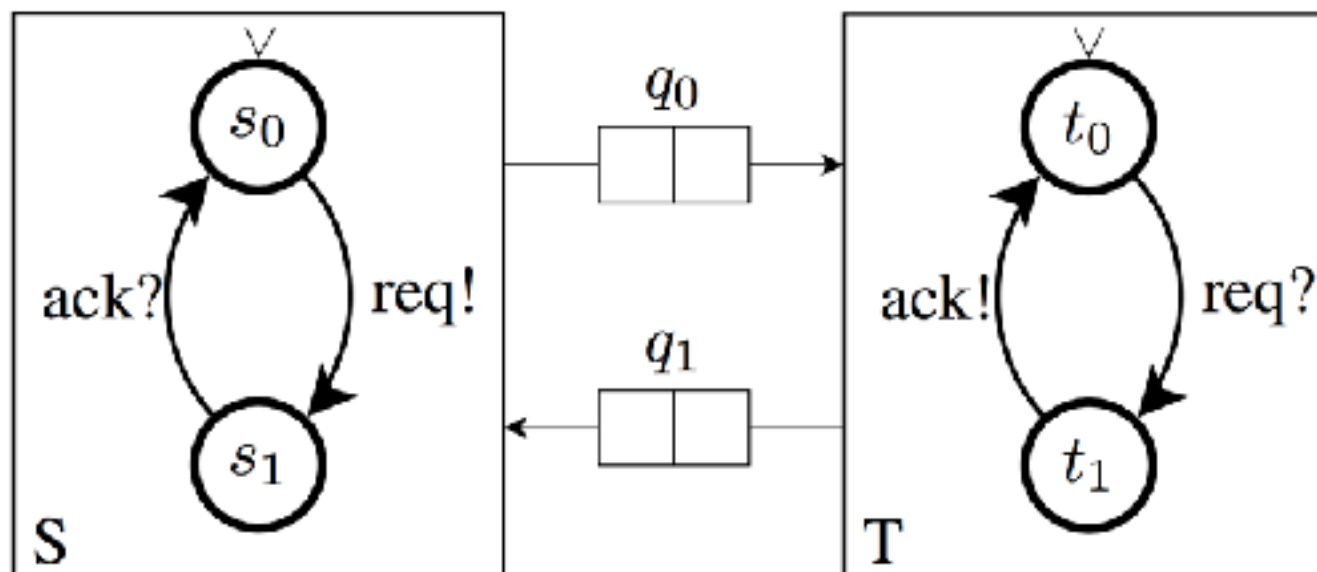
- $\# \text{A} = \# \text{B}$



## Invariant Generation Highlight (2)

Cross-layer invariants:

- Automatically generated
- Linear equations over:
  - (1) # transitions and being in a given state
  - (2) # of “events” on a channel and # transitions

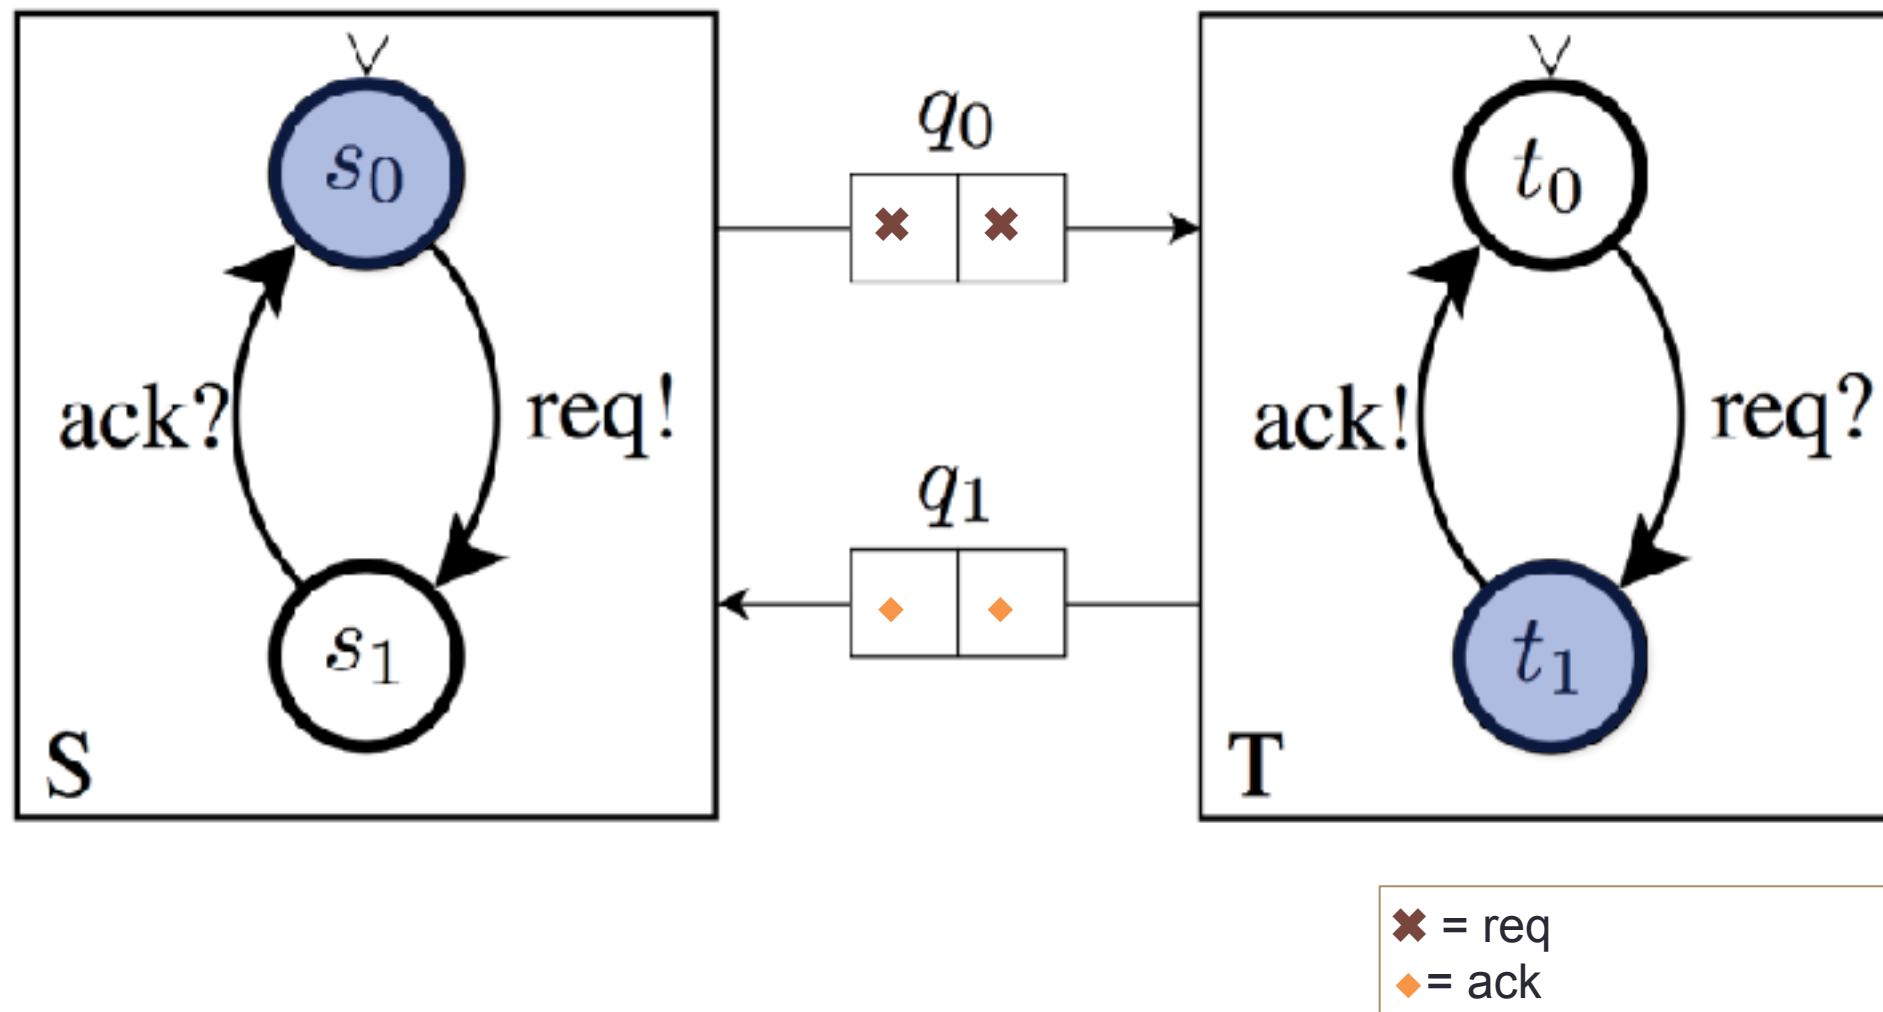


$$T.t_0 - S.s_0 = \#q_0.req + \#q_1.ack$$





## Invariant Generation Highlight (2)



$$T.t_0 - S.s_0 = \#q_0.req + \#q_1.ack$$

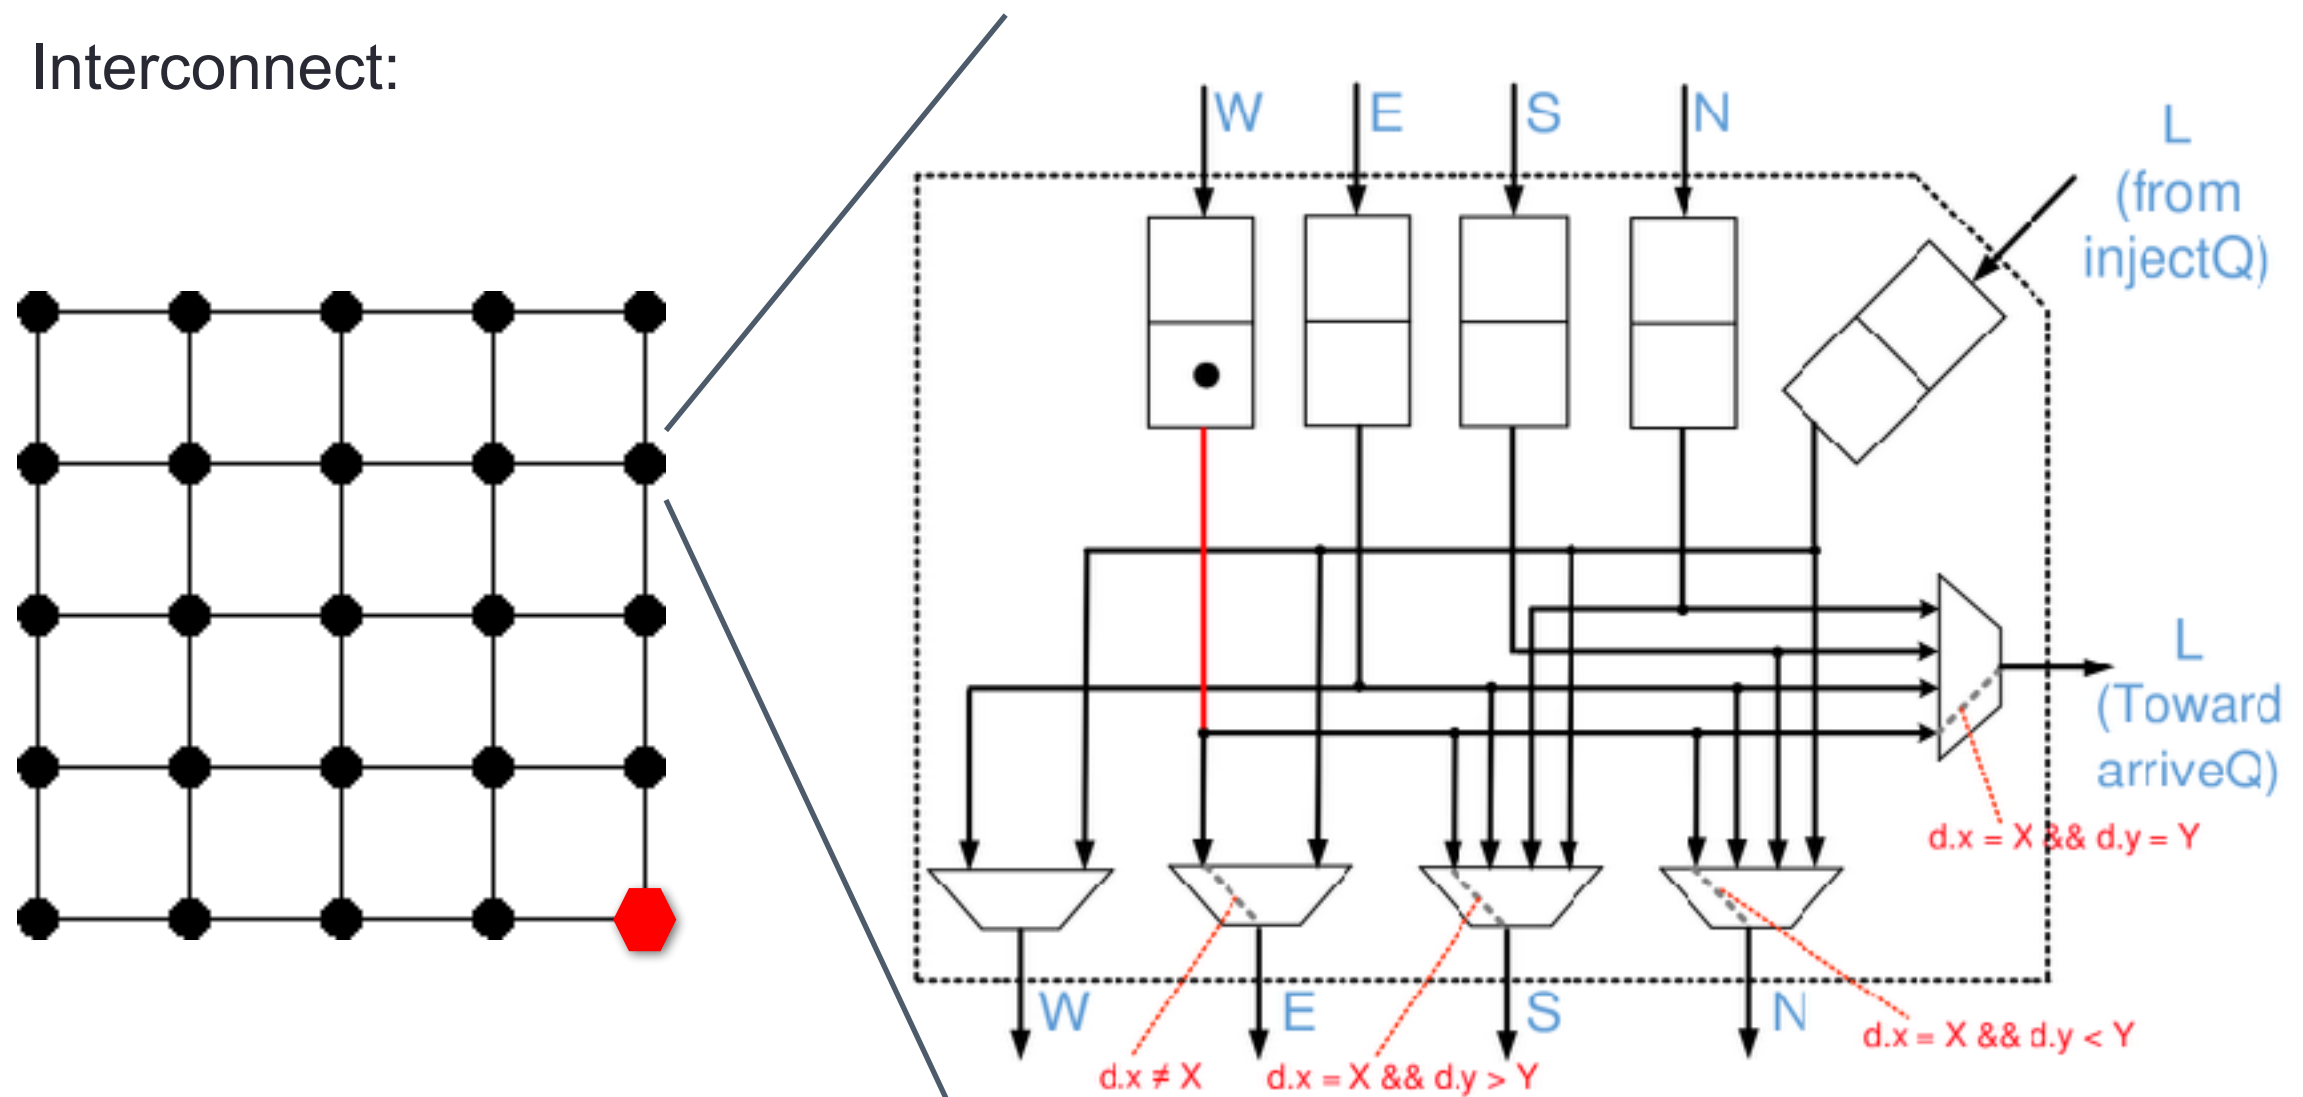


# Experimental results



# Case-study: 2D Mesh - XY routing - MI protocol

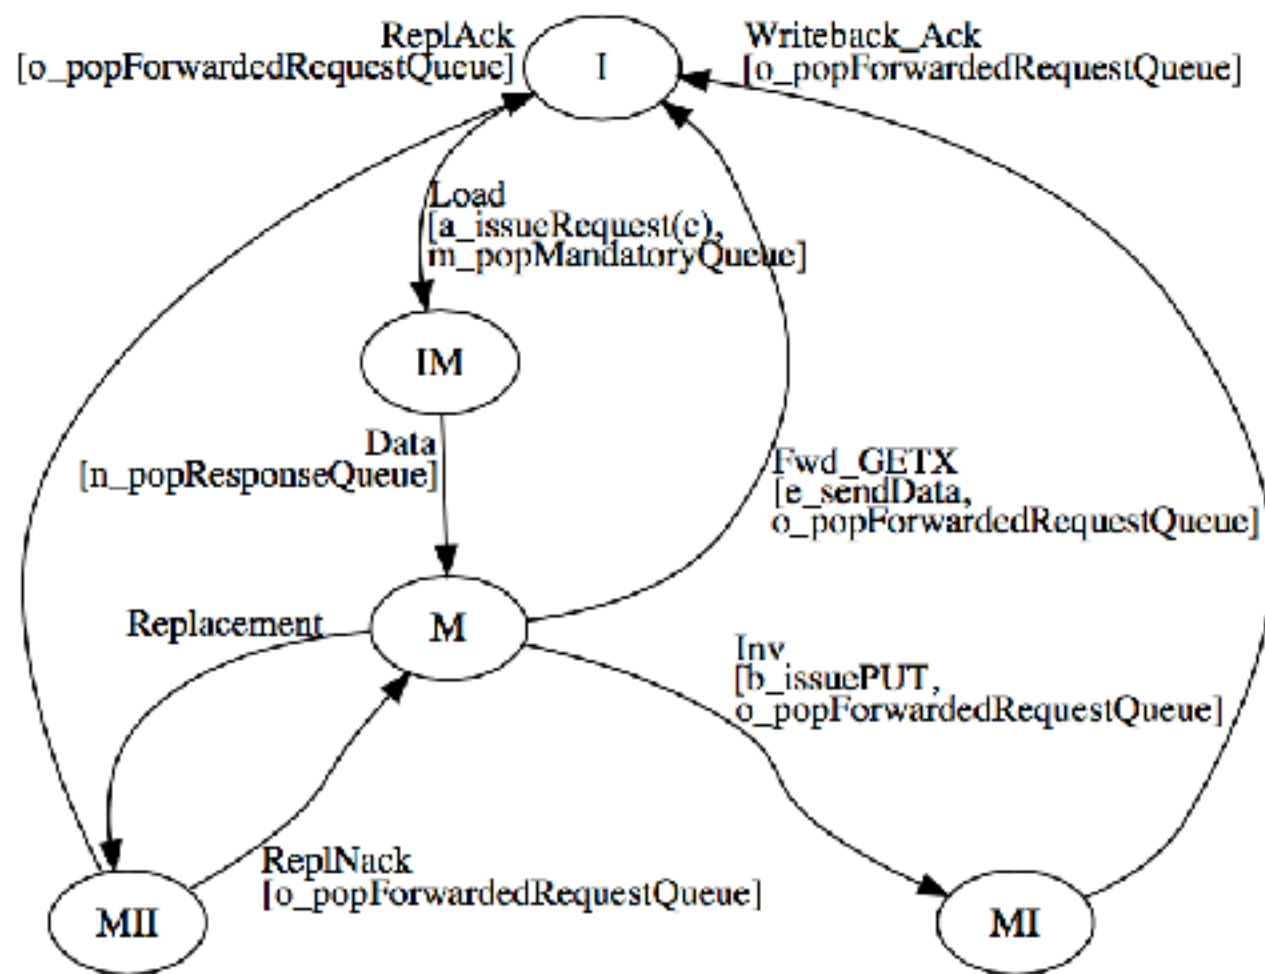
Interconnect:



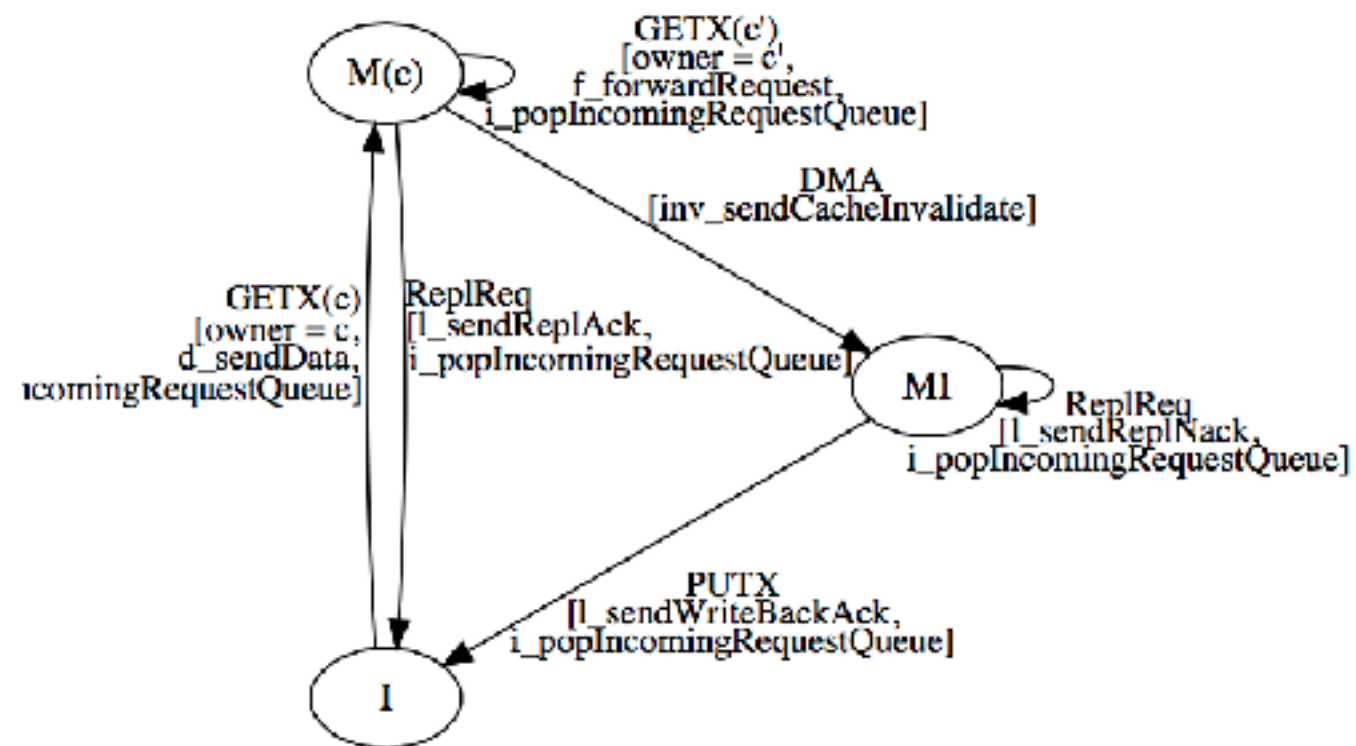


# Case-study: 2D Mesh - XY routing - MI protocol

Protocol:



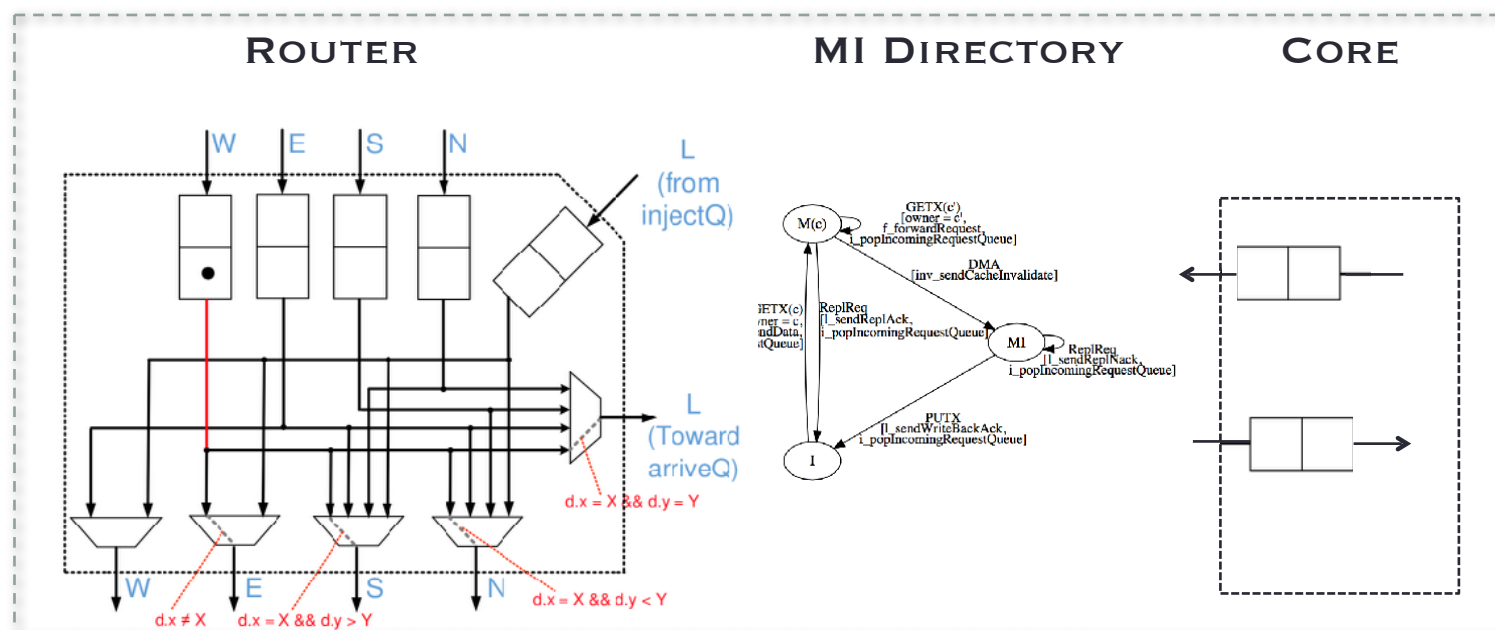
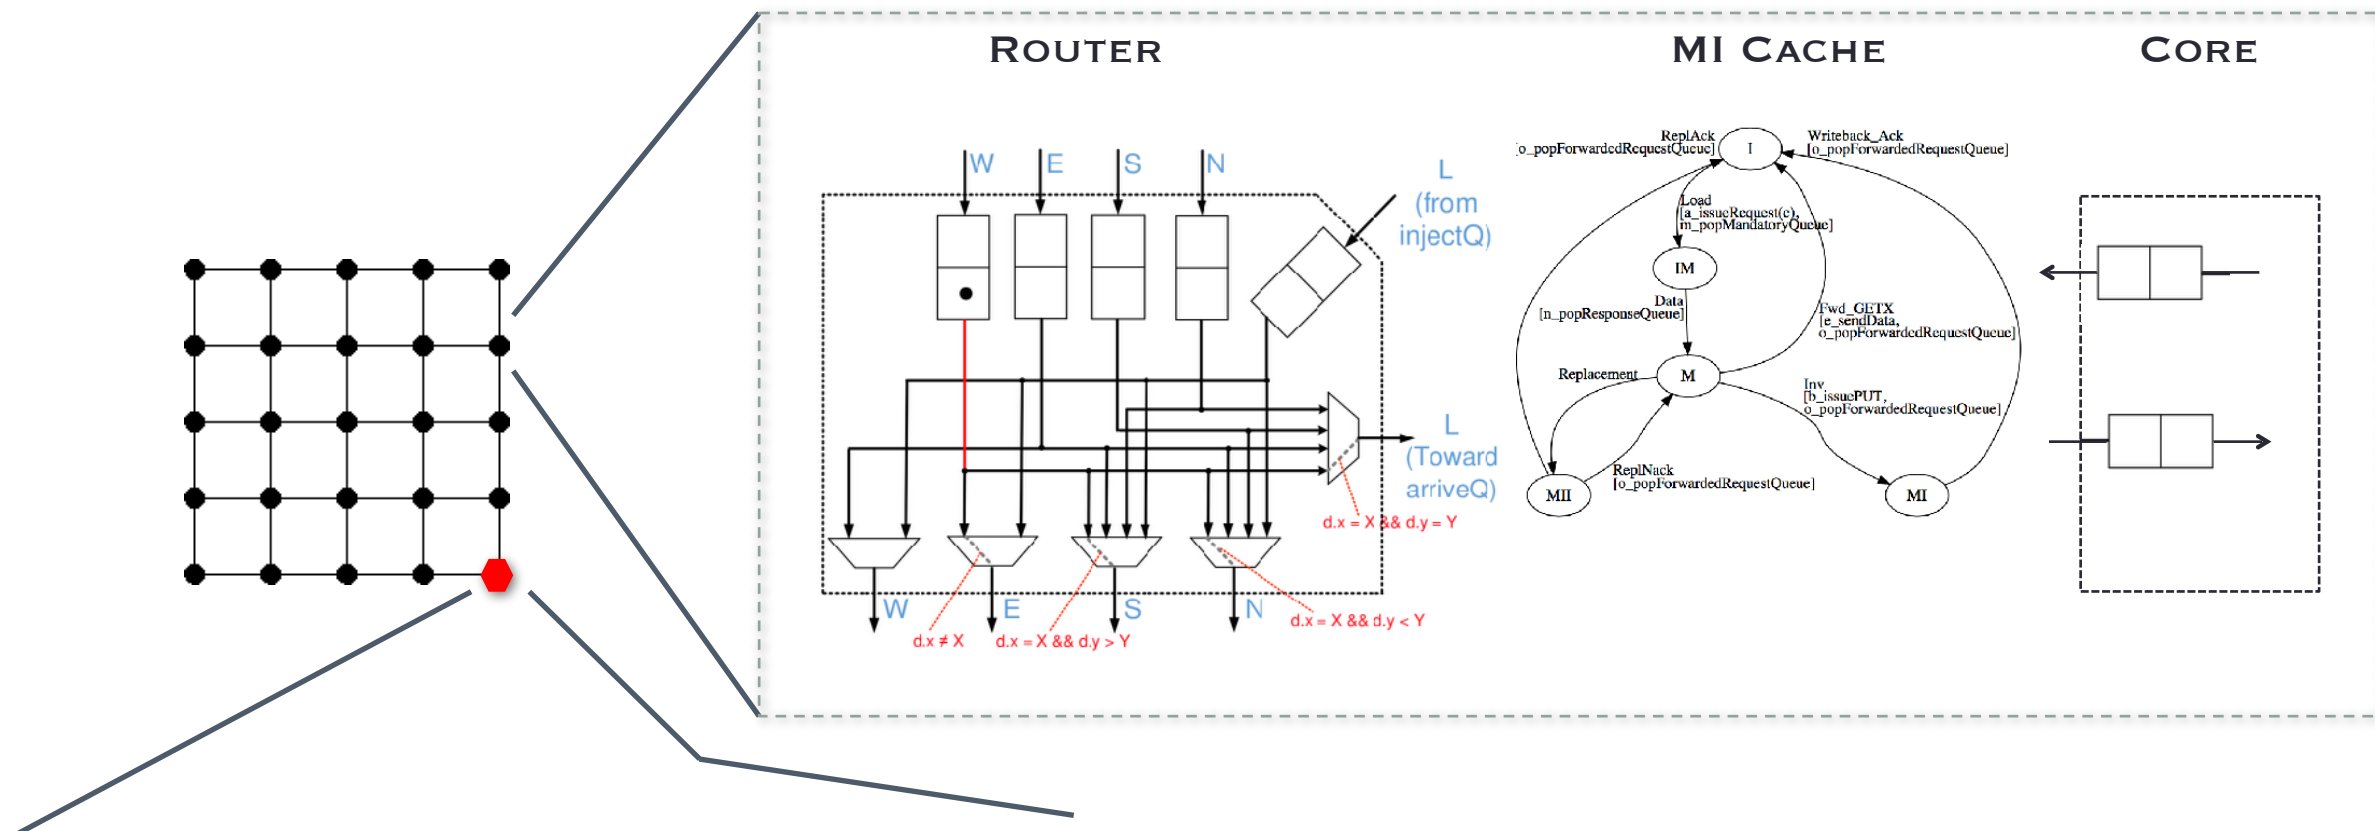
L2 caches



Directory



# Case-study: 2D Mesh - XY routing - MI protocol





## Case-study: Experimental results


<b>Size</b>	<b>DL</b>	<b>DLF</b>	<b>#primitives</b>	<b>#queues</b>	<b>#automata</b>
$2 \times 2$	1.7s	1.3s	100	24	4
$3 \times 3$	23s	16s	225	54	9
$4 \times 4$	3m52s	2m41s	400	96	16
$5 \times 5$	33m18	23m5s	625	150	25



# Reachability analysis



# MaDL Verification: Overview



Generate Invariants

Generate SMT Problem

SMT  
Solver (Z3)

(unsat)



No deadlock!

(sat)



Generate SMV Model

nuXmv



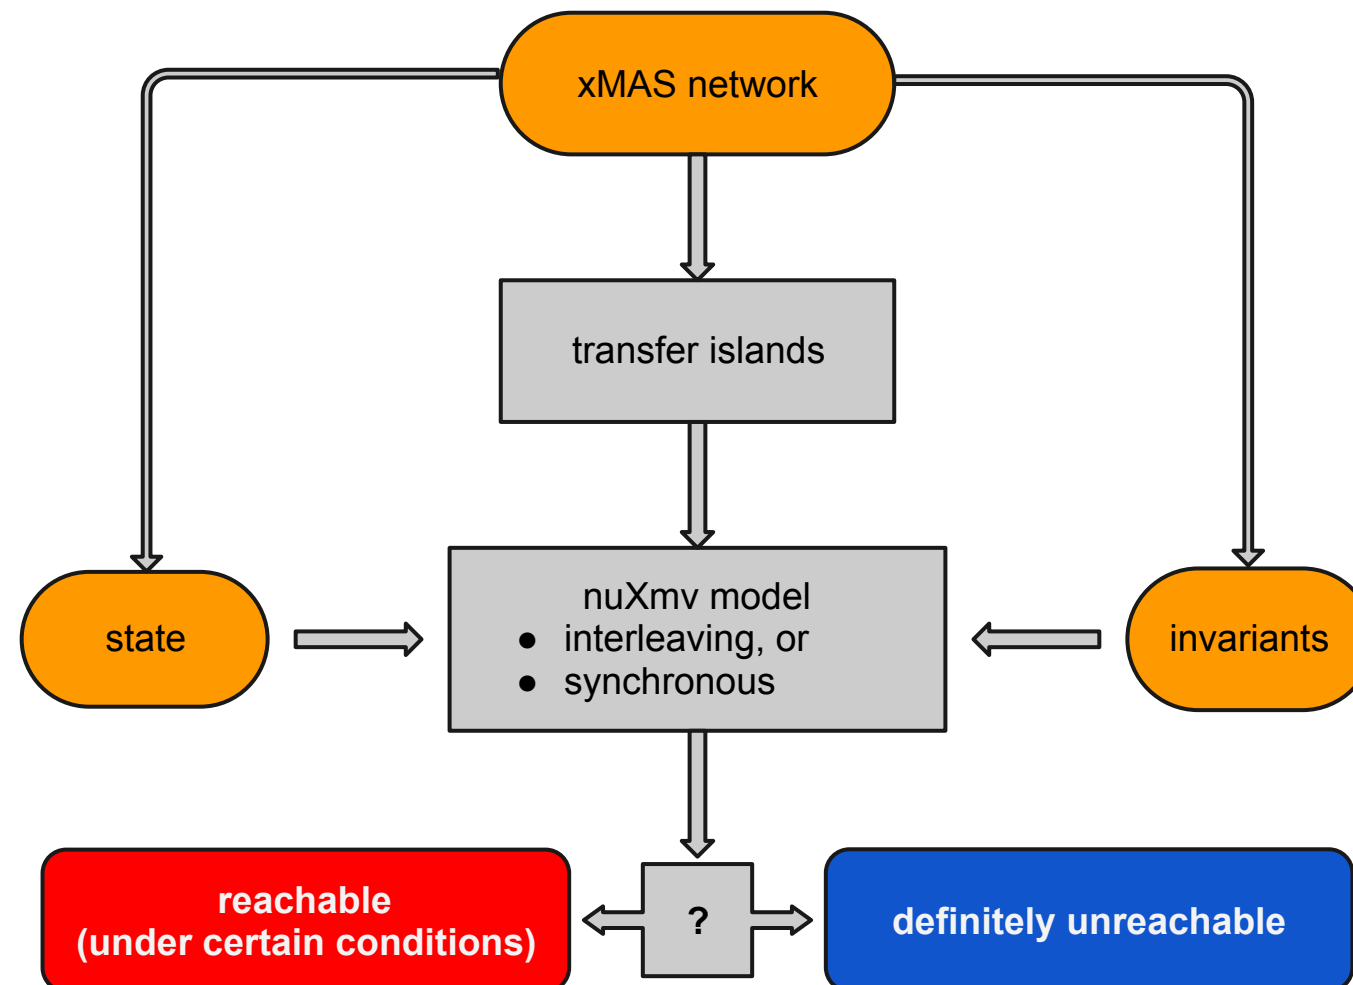
Deadlock!

Basic idea:

- Express deadlock in SMT
- Over-approximate deadlocks
- Use invariants to rule-out false deadlocks
- Reachability analysis of found deadlocks




# Reachability checking flow




# Transfer islands

---




# Enabled island

---




# Enabled island

---




# Combination of islands

---




# Combination of islands

---



# Conflicts

---



# Different semantics

---

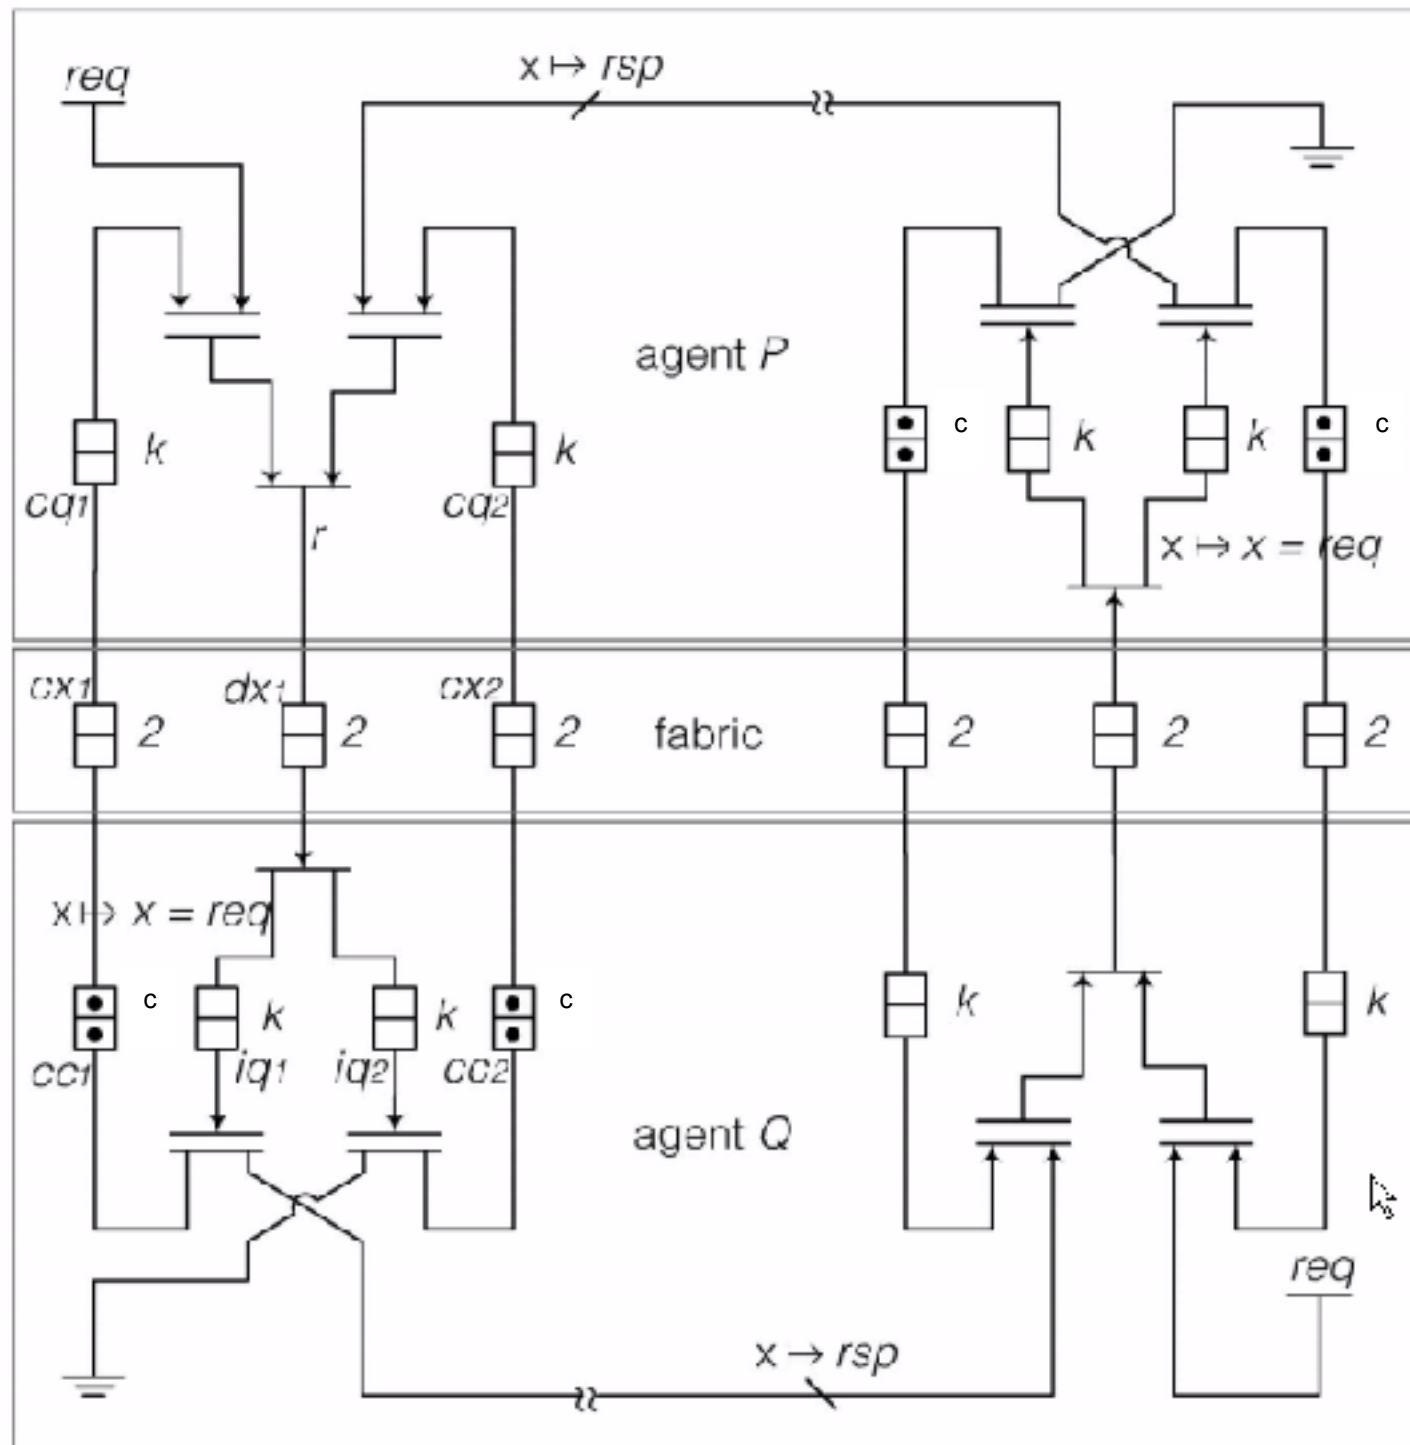
*interleaving*

*asynchronous*

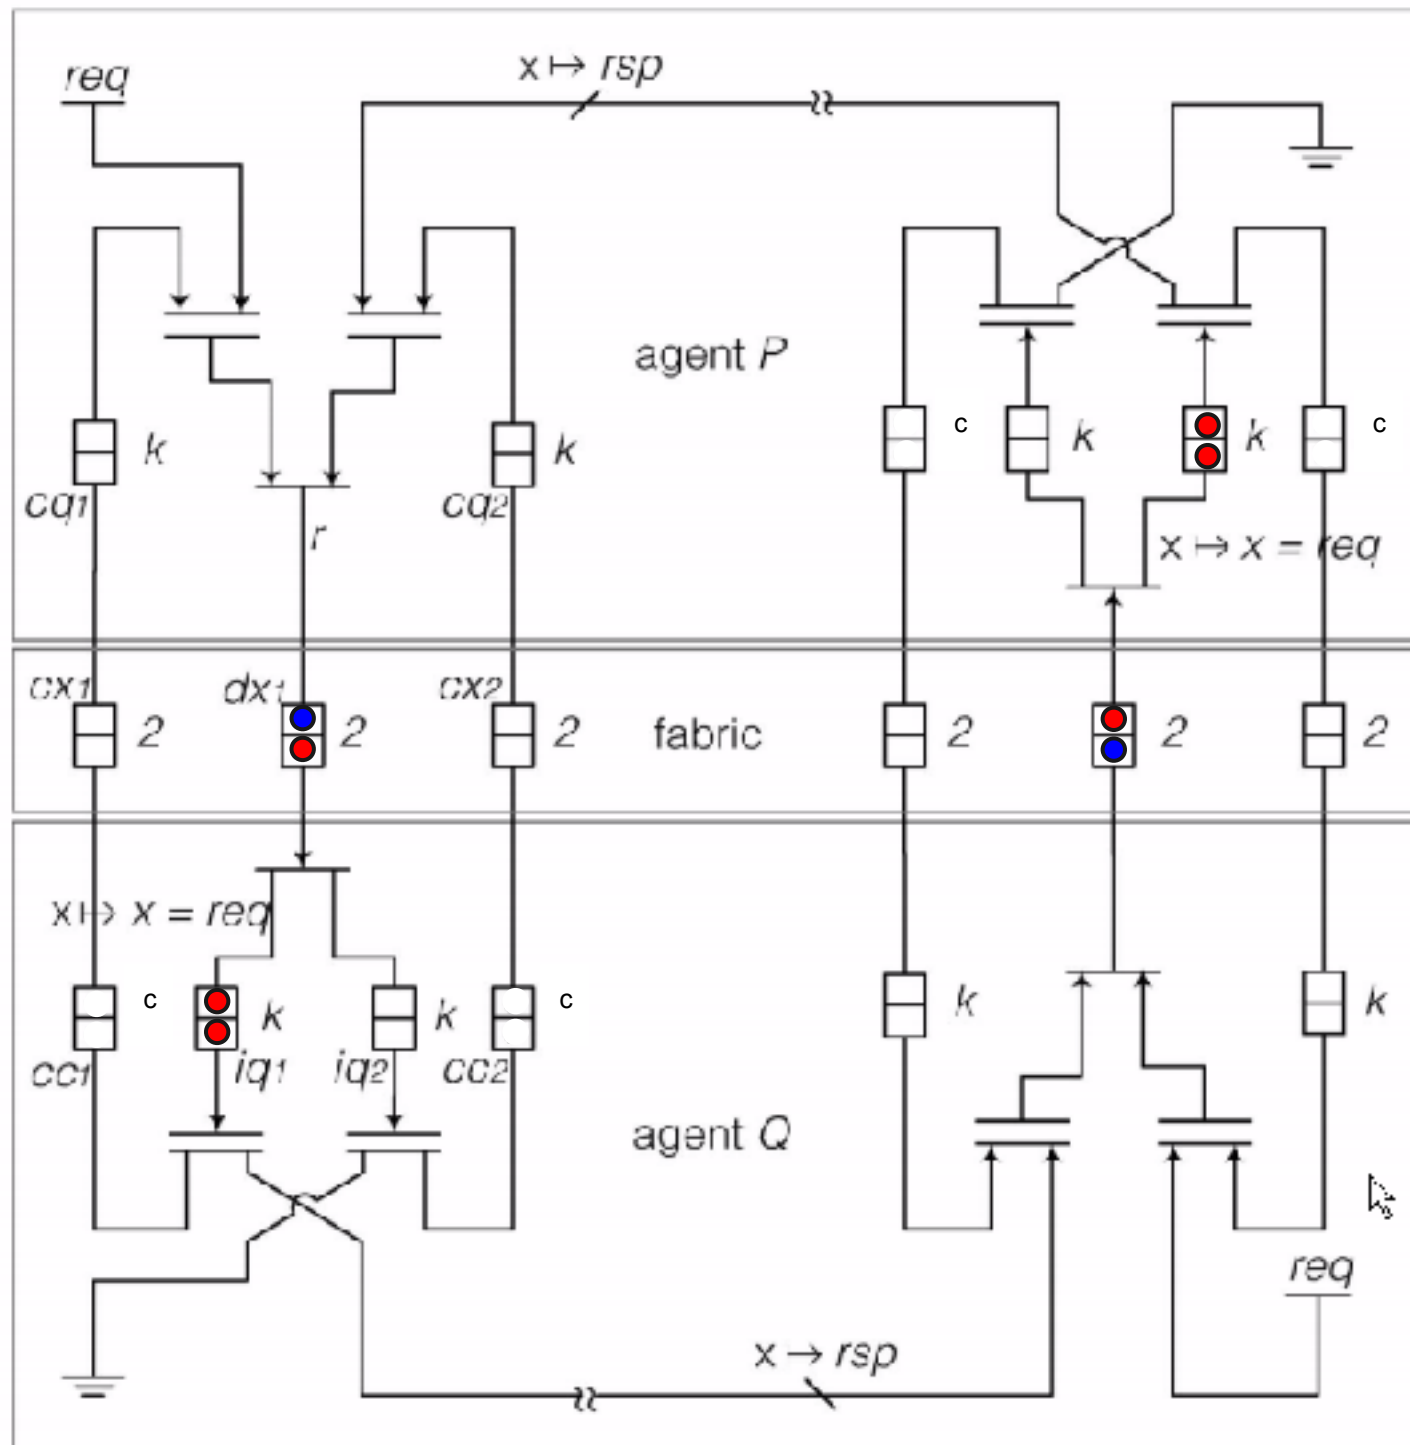
*synchronous*



# Two Agents Example



# Deadlock found by SMT

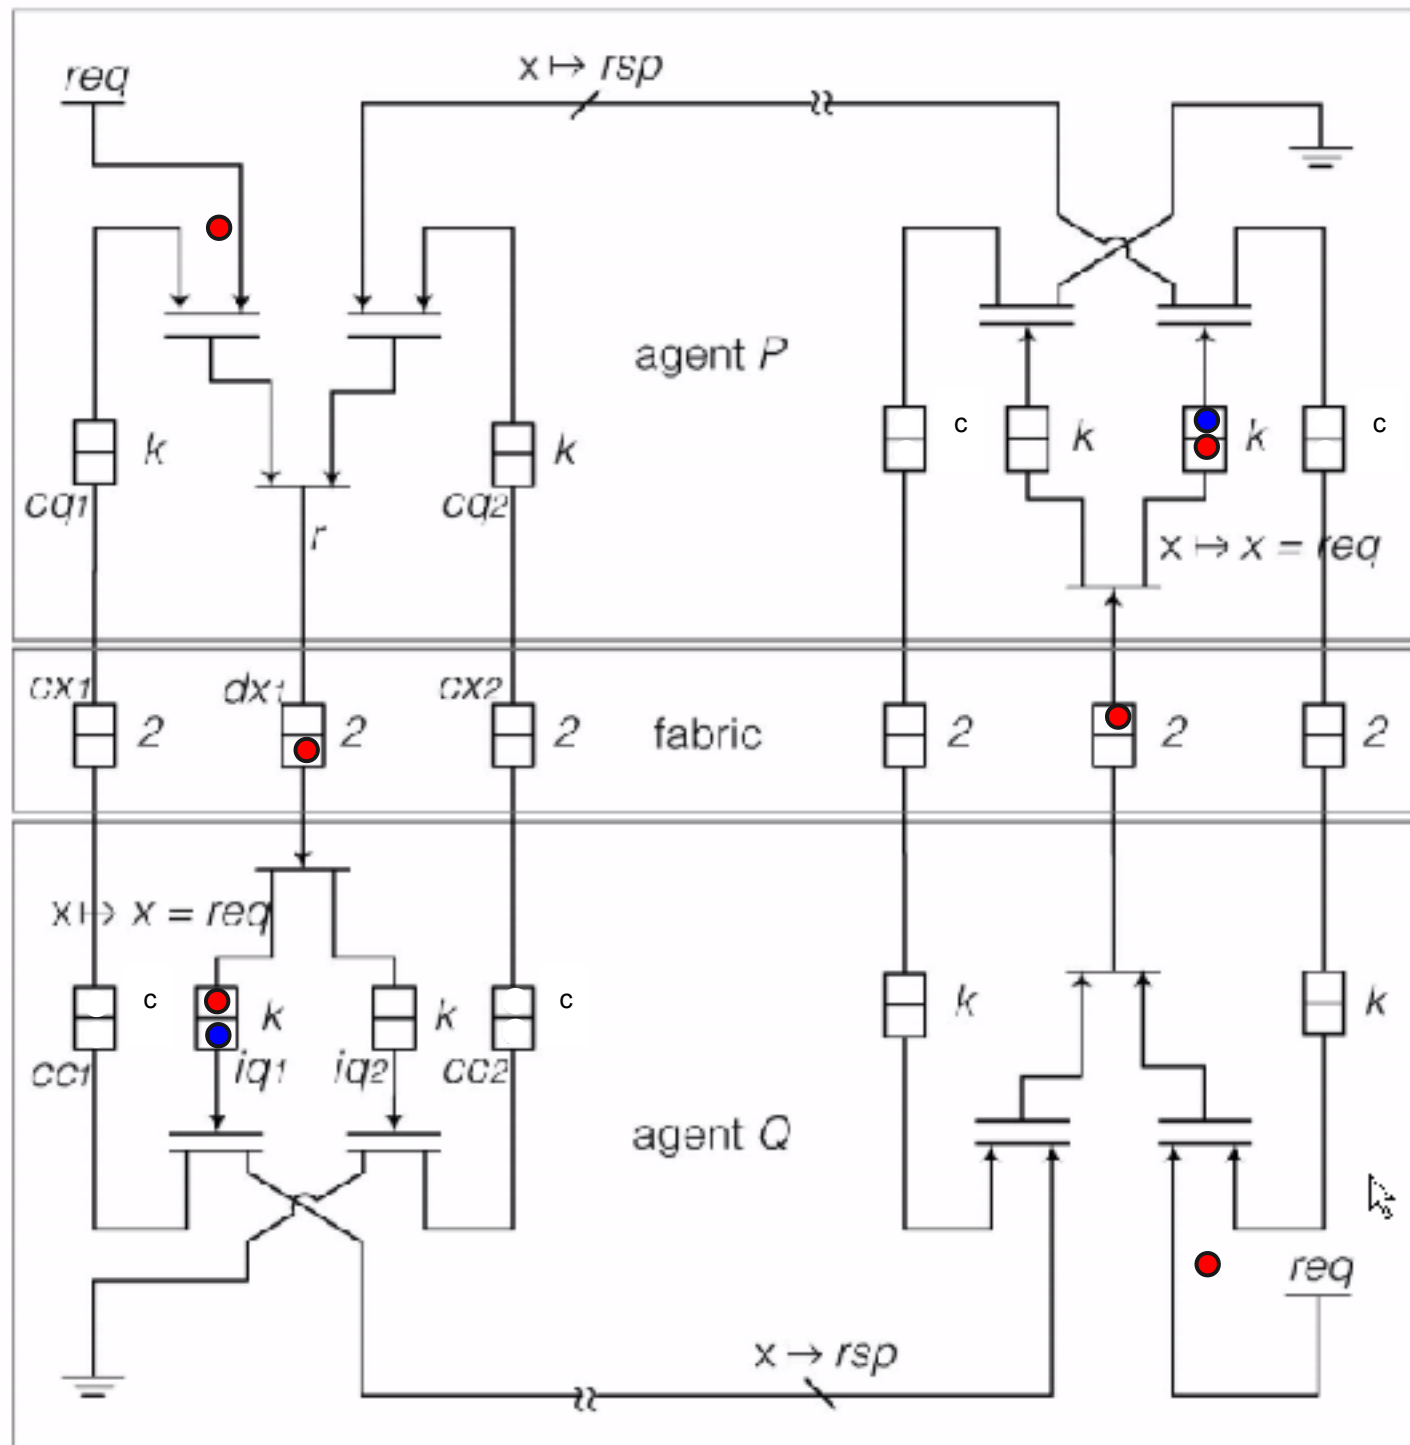


"However if credit counters are sized incorrectly to provide more credits (say  $k + 1$ ) than the capacity  $k$  of the ingress queues, the system deadlocks ..."

-- Chatterjee et al. VMCAI 2011

**Incorrect!**

# Let's rewind one cycle



No credit available.  
(3 requests)



# Some experimental results

---

- Two agents example
- IC3 or BDD
- optional invariants
- synchronous or interleaving

<b>c = 2; k = 2</b>	unreachable deadlock
<b>2; 3</b>	<b>unreachable deadlock</b>
<b>4; 4</b>	unreachable deadlock
<b>4; 5</b>	<b>unreachable deadlock</b>
<b>4; 6</b>	reachable deadlock

synchronous/IC3+INV ~3s

state space:  $\sim 2^{40}$  (about  $10^{12}$ )

# Outline

---

- » Intel's micro-architectural description language
  - xMAS language
  - Capturing high-level structure and message dependencies
  - Extended to “MaDL” at TU/e.
    - Micro-architectural Description Language
- » Deadlock verification for MaDL
  - Definition of deadlocks
  - Labelled dependency graph
  - Feasible logically closed subgraph
- » **Conclusion and future work**

# Research questions and axes

**Q1: How can we formalise the generic aspects of specific domains?**

Foundations

Formalise domain-specific theories

**Q2: How can we define specification languages with built-in support for formal analyses?**

Abstractions

Focus on aspects and properties

Language to express and analyse

**Q3: How do we relate the formally proven correct specification to the actual design?**

Approximations

Ignore details to scale-up

# Research questions

- Q1: How can we formalise the generic aspects of specific domains?
- Q2: How can we define specification languages with built-in support for formal analyses?
- Q3: How do we relate the formal proven correct specification to the actual design?



**General Theory of Networking Architectures (Q1)**

GeNoC theory with Productivity Theorem  
Deadlock-free routing theory & algorithms

**Languages (Q2)**

**Verification  
Technologies  
(Q2 & Q3)**

MaDL language & algorithms

**Integration**

Cooperation with industrial partners  
(Intel,ARM)

**Real World Applications**

and academic partners  
(FBK, UC Irvine, Tsinghua, ...)

# Today's focus

## General Theory of Networking Architectures (Q1)

GeNoC theory with Productivity Theorem

Deadlock-free routing theory & algorithms

## Languages (Q2)

## Verification Technologies (Q2 & Q3)

MaDL language & algorithms

## Integration

Cooperation with industrial partners  
(Intel, ARM)

## Real World Applications

and academic partners  
(FBK, UC Irvine, Tsinghua, ...)



# MaDL: Conclusions & Future Work

- Conclusions
  - Adequate DSL for prototyping network architectures.
  - Being applied to some industrial case-studies
- Future work
  - Equivalence relations between MaDL and RTL
  - Performance (throughput & latency) evaluation
  - Extend to entire systems



# MaDL: Current activities

- Master theses
  - Perry: analysis for ring networks
  - Ruud: Quality-of-Services
- PhD theses
  - Alexander: equivalences, pre-orders, link with RTL
- Valorisation
  - Precuneus Solutions B.V.

# MaDL Whiteboard DEMO



THANKS!

